

A preface: Everything is done in 10 or 20 epochs because of the lack of time. I apologize for that in advance.

## Task 1

```
In [ ]: import tensorflow as tf
        from tensorflow import keras
        from keras import layers
        import numpy as np
        import matplotlib.pyplot as plt

        batch_size = 8

        # Build a simple CNN with strided convolution layers
        def define_model():
            inputs = tf.keras.Input(shape=(28,28,1),name='Inputs')
            x = layers.Conv2D(16,kernel_size=(5,5),activation='relu',padding='same',strides=1,name='L1')(inputs)
            x = layers.Conv2D(16,kernel_size=(3,3),activation='relu',padding='same',strides=2,name='L2')(x)
            x = layers.Conv2D(16,kernel_size=(3,3),activation='relu',padding='same',strides=1,name='L3')(x)
            x = layers.Conv2D(16,kernel_size=(3,3),activation='relu',padding='same',strides=2,name='L4')(x)
            x = layers.Flatten()(x)
            embedding_layer = layers.Dense(2, name='Embedding')(x)
            outputs = layers.Dense(10,activation='softmax')(embedding_layer)
            model = keras.Model(inputs=inputs, outputs=outputs)
            return model

        def mnist_data():
            (x_train, y_train), (x_test, y_test) = keras.datasets.mnist.load_data()
            y_train = tf.keras.utils.to_categorical(y_train)
            y_test = tf.keras.utils.to_categorical(y_test)
            x_train = x_train/255.0
            x_test = x_test/255.0
            return x_train, y_train, x_test, y_test

        # Train the model on MNIST data using standard cross-entropy loss
```

```
def train_model(model, x_train, y_train, x_test, y_test, epochs=10):
    model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
    model.fit(x_train, y_train, validation_data=(x_test, y_test), epochs=epochs, batch_size=batch_size)

x_train, y_train, x_test, y_test = mnist_data()
model = define_model()
train_model(model, x_train, y_train, x_test, y_test, 10)
```

Epoch 1/10

7500/7500 [=====] - 39s 5ms/step - loss: 0.6888 - accuracy: 0.7824 - val\_loss: 0.3921 - val\_accuracy: 0.8943

Epoch 2/10

7500/7500 [=====] - 38s 5ms/step - loss: 0.3534 - accuracy: 0.9089 - val\_loss: 0.2811 - val\_accuracy: 0.9314

Epoch 3/10

7500/7500 [=====] - 38s 5ms/step - loss: 0.2685 - accuracy: 0.9320 - val\_loss: 0.2456 - val\_accuracy: 0.9388

Epoch 4/10

7500/7500 [=====] - 40s 5ms/step - loss: 0.2238 - accuracy: 0.9430 - val\_loss: 0.2062 - val\_accuracy: 0.9509

Epoch 5/10

7500/7500 [=====] - 42s 6ms/step - loss: 0.1964 - accuracy: 0.9506 - val\_loss: 0.1956 - val\_accuracy: 0.9538

Epoch 6/10

7500/7500 [=====] - 39s 5ms/step - loss: 0.1771 - accuracy: 0.9546 - val\_loss: 0.1828 - val\_accuracy: 0.9566

Epoch 7/10

7500/7500 [=====] - 39s 5ms/step - loss: 0.1619 - accuracy: 0.9594 - val\_loss: 0.1924 - val\_accuracy: 0.9561

Epoch 8/10

7500/7500 [=====] - 39s 5ms/step - loss: 0.1521 - accuracy: 0.9610 - val\_loss: 0.1737 - val\_accuracy: 0.9599

Epoch 9/10

7500/7500 [=====] - 40s 5ms/step - loss: 0.1437 - accuracy: 0.9632 - val\_loss: 0.1703 - val\_accuracy: 0.9608

Epoch 10/10

7500/7500 [=====] - 40s 5ms/step - loss: 0.1371 - accuracy: 0.9643 - val\_loss: 0.1735 - val\_accuracy: 0.9608

```
In [ ]: model_for_embeddings = tf.keras.Model(inputs=model.input,
                                              outputs=model.get_layer('Embedding').output)
```

```

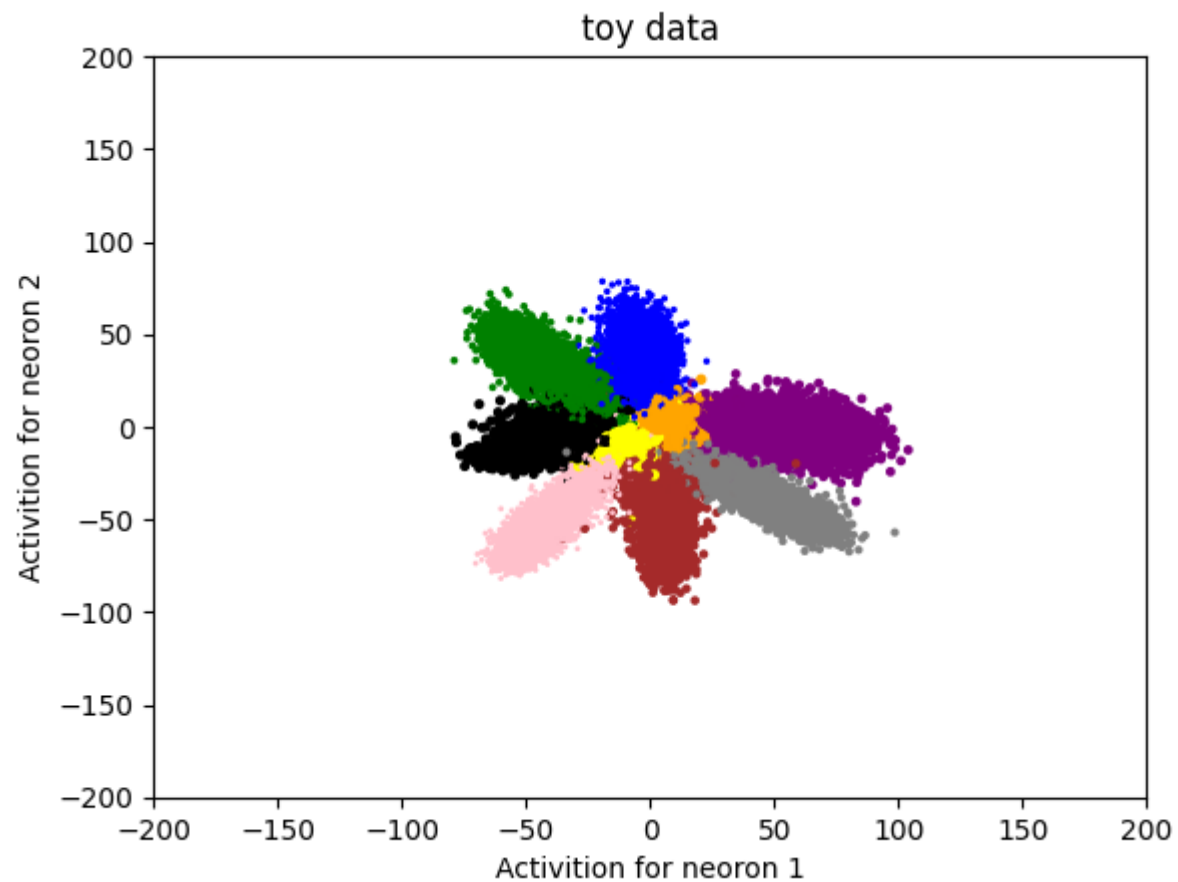
In [ ]: labels = y_train.argmax(axis=1)
        outs = model_for_embeddings.predict(x_train)

        colors = {
            0: 'red',
            1: 'pink',
            2: 'blue',
            3: 'green',
            4: 'grey',
            5: 'brown',
            6: 'purple',
            7: 'black',
            8: 'orange',
            9: 'yellow',
        }

        %matplotlib inline
        fig, ax = plt.subplots()
        ax.scatter(outs[:,0], outs[:,1], labels, c=[colors[x] for x in labels])
        ax.set_xlabel('Activition for neuron 1')
        ax.set_ylabel('Activition for neuron 2')
        ax.set_title('toy data')
        ax.set_xlim(-200,200)
        ax.set_ylim(-200,200)
        plt.show()

```

1875/1875 [=====] - 6s 3ms/step



## Task 2

```
In [ ]: import tensorflow as tf
        from tensorflow import keras
        from keras import layers
        import numpy as np
        import matplotlib.pyplot as plt

        batch_size = 8
        alpha = 1
        learning_rate = 0.0001
```

```

@tf.function
def contrastive_loss(y_true, y_pred):
    y_true = tf.argmax(y_true, axis=1)
    L=0
    for i in range(batch_size):
        for j in range(i+1, batch_size):
            D = tf.square(tf.norm(y_pred[i,:] - y_pred[j,:]))
            if y_true[i] == y_true[j]:
                L = L+D
            else:
                L = L+tf.maximum(0.0, alpha-D)
    return L

def mnist_data():
    (x_train, y_train), (x_test, y_test) = keras.datasets.mnist.load_data()
    y_train = tf.keras.utils.to_categorical(y_train)
    y_test = tf.keras.utils.to_categorical(y_test)
    x_train = x_train/255.0
    x_test = x_test/255.0
    return x_train, y_train, x_test, y_test

# Build a simple CNN with strided convolution layers
def define_model2():
    inputs = tf.keras.Input(shape=(28,28,1),name='Inputs')
    x = layers.Conv2D(16,kernel_size=(5,5),activation='relu',padding='same',strides=1,name='L1')(inputs)
    x = layers.Conv2D(16,kernel_size=(3,3),activation='relu',padding='same',strides=2,name='L2')(x)
    x = layers.Conv2D(16,kernel_size=(3,3),activation='relu',padding='same',strides=1,name='L3')(x)
    x = layers.Conv2D(16,kernel_size=(3,3),activation='relu',padding='same',strides=2,name='L4')(x)
    x = layers.Flatten()(x)
    x = layers.Dense(2, name='Embedding')(x)
    x = layers.Dense(10,activation='softmax')(x)
    model = keras.Model(inputs=inputs, outputs=x)
    return model

x_train, y_train, x_test, y_test = mnist_data()

# Train the model on MNIST data using standard cross-entropy loss
def train_model2(model, x_train, y_train, x_test, y_test, epochs=100):
    model.compile(optimizer=tf.keras.optimizers.Adam(), loss=contrastive_loss, metrics=['accuracy'])

```

```
model.fit(x_train, y_train, validation_data=(x_test, y_test), epochs=epochs, batch_size=batch_size)
```

```
model2 = define_model2()  
train_model2(model2, x_train, y_train, x_test, y_test, 10)
```

Epoch 1/10

7500/7500 [=====] - 52s 6ms/step - loss: 5.2253 - accuracy: 0.1699 - val\_loss: 3.2521 - val\_accuracy: 0.1875

Epoch 2/10

7500/7500 [=====] - 47s 6ms/step - loss: 3.0373 - accuracy: 0.1861 - val\_loss: 2.4646 - val\_accuracy: 0.1774

Epoch 3/10

7500/7500 [=====] - 47s 6ms/step - loss: 2.4444 - accuracy: 0.1899 - val\_loss: 1.9718 - val\_accuracy: 0.1882

Epoch 4/10

7500/7500 [=====] - 48s 6ms/step - loss: 2.1174 - accuracy: 0.1932 - val\_loss: 1.7910 - val\_accuracy: 0.1856

Epoch 5/10

7500/7500 [=====] - 48s 6ms/step - loss: 1.9985 - accuracy: 0.1933 - val\_loss: 1.8336 - val\_accuracy: 0.1830

Epoch 6/10

7500/7500 [=====] - 49s 7ms/step - loss: 1.8747 - accuracy: 0.1972 - val\_loss: 1.8118 - val\_accuracy: 0.1987

Epoch 7/10

7500/7500 [=====] - 48s 6ms/step - loss: 1.7875 - accuracy: 0.2126 - val\_loss: 1.3228 - val\_accuracy: 0.2683

Epoch 8/10

7500/7500 [=====] - 48s 6ms/step - loss: 1.3673 - accuracy: 0.2737 - val\_loss: 1.1645 - val\_accuracy: 0.2687

Epoch 9/10

7500/7500 [=====] - 48s 6ms/step - loss: 1.2331 - accuracy: 0.2775 - val\_loss: 1.0885 - val\_accuracy: 0.2753

Epoch 10/10


7500/7500 [=====] - 46s 6ms/step - loss: 1.1842 - accuracy: 0.2802 - val\_loss: 1.1453 - val\_accuracy: 0.2690

```
In [ ]: labels = y_train.argmax(axis=1)  
        model2_for_embeddings = tf.keras.Model(inputs=model2.input,  
                                                outputs=model2.get_layer('Embedding').output)  
        outs = model2_for_embeddings.predict(x_train)
```

```
colors = {
    0: 'red',
    1: 'pink',
    2: 'blue',
    3: 'green',
    4: 'grey',
    5: 'brown',
    6: 'purple',
    7: 'black',
    8: 'orange',
    9: 'yellow',
}

%matplotlib inline
fig, ax = plt.subplots()
ax.scatter(outs[:,0], outs[:,1], labels, c=[colors[x] for x in labels])
ax.set_xlabel('Activition for neuron 1')
ax.set_ylabel('Activition for neuron 2')
ax.set_title('toy data')
plt.show()
```

1875/1875 [=====] - 6s 3ms/step



In [ ]: