

Image Segmentation

Advanced Computer Vision


Niels Landwehr

Overview

- Introduction: Computer Vision
- Data, Models, Optimization
- Neural Networks and Automatic Differentiation
- Convolutional Architectures For Image Classification
- Visualization and Transfer Learning
- Metric Learning
- **Image Segmentation**

Problem Setting Segmentation

- Problem setting image segmentation
 - Input: image as $m \times l \times d$ tensor (typically, $d=3$).
 - Output: for each of the pixels, a class label from a predefined set of classes $\{c_1, \dots, c_k\}$ (semantic category).
 - Typically, output is an $m \times l \times k$ tensor of class scores.



{Trees, Sky,
Cow, Grass, ...}


- More fine-grained analysis compared to classification or object detection.
- **To arrive at segmentation, model needs to (implicitly) identify and classify objects in image and find their exact contours**

Problem Setting Segmentation

- Many applications: autonomous driving, medical domains, ...



Cordts, Marius, et al.
"The cityscapes dataset."
*CVPR Workshop on the Future
of Datasets in Vision*. Vol. 2.
2015.




Brain tumor segmentation in MRI images

Chen, Chen, et al. "3D dilated multi-fiber network for real-time brain tumor segmentation in MRI." *International Conference on Medical Image Computing and Computer-Assisted Intervention*. Springer, Cham, 2019.

Network Architectures for Segmentation?

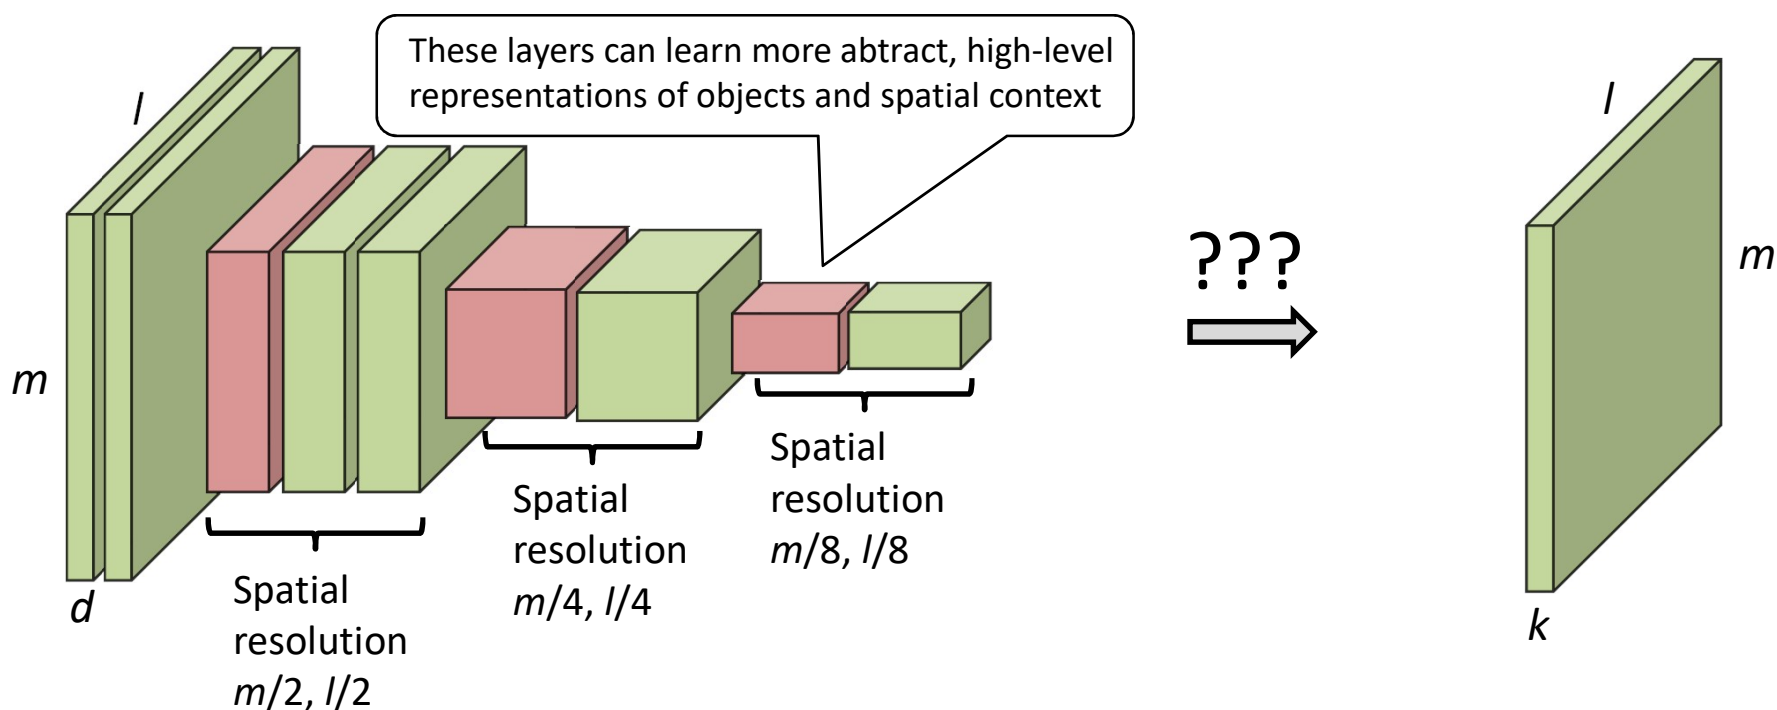
- How can we design a network architecture that maps $m \times l \times d$ tensors to $m \times l \times k$ tensors (and implicitly reasons about spatial objects)?
- First idea: use convolution layers with spatial dimension $m \times l$ throughout



- Problem: without spatial pooling, difficult to learn about larger-scale objects in images. Without large-scale context cannot perform good segmentation.
- Problem: convolutions at full spatial resolutions are expensive.

Spatial Downsampling

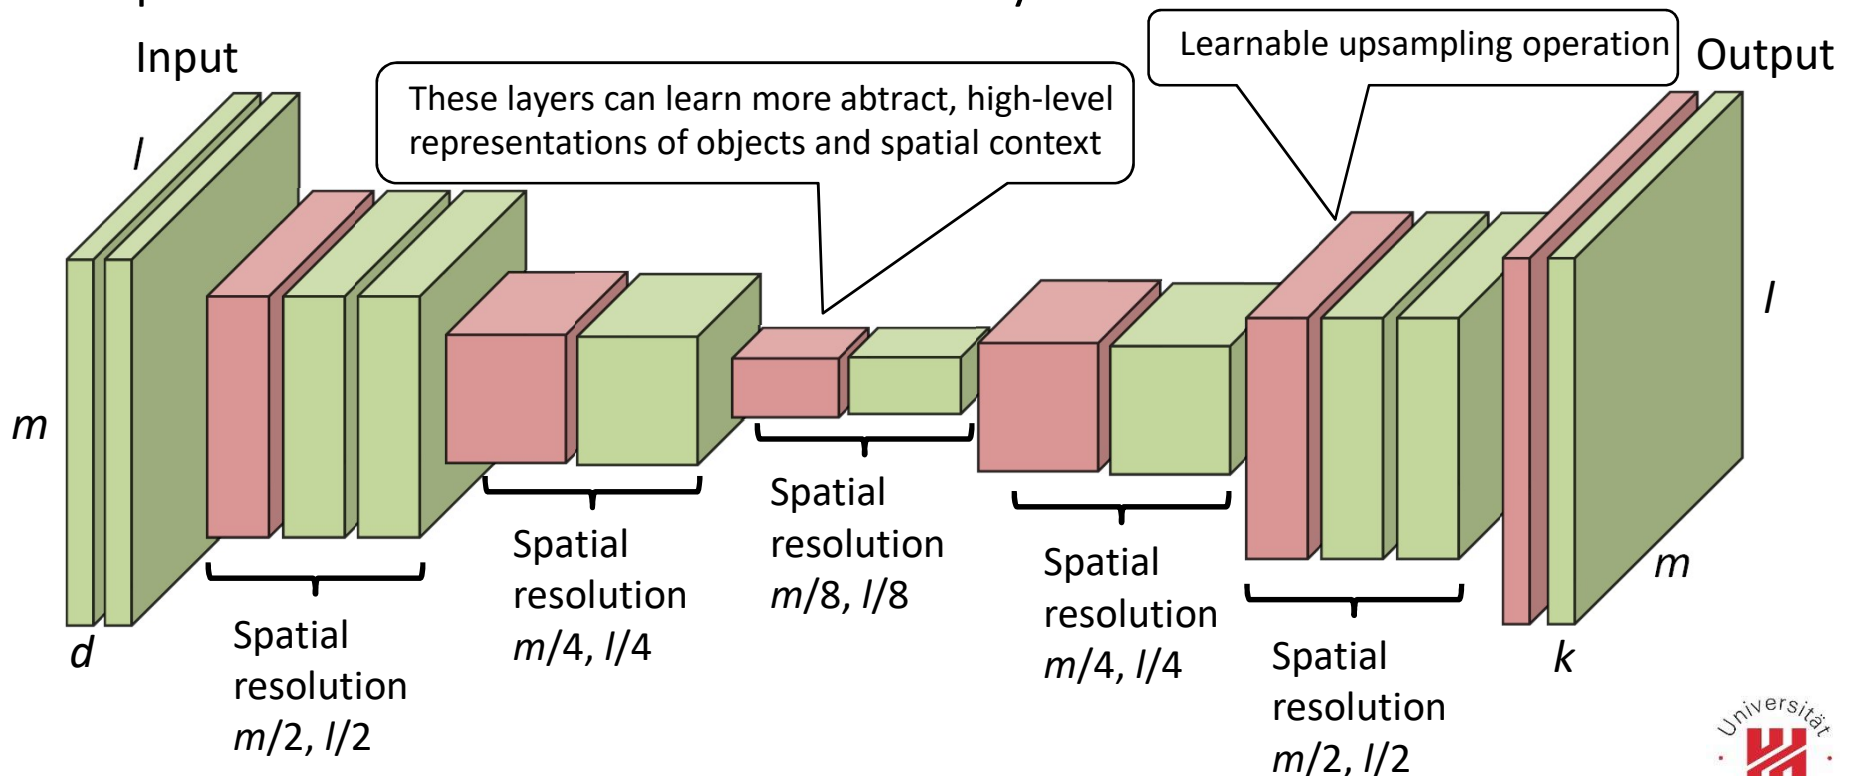
- To reason about large-scale objects (and reduce computational cost), should use pyramid-shaped architecture as in classification networks
- Pooling or strided convolution layers to reduce spatial dimension



- However, need high resolution in output. How do we go back?

Spatial Downsampling and Upsampling

- **Idea:** first downsample by pooling and strided convolutions, then upsample again to the final resolution
- Learnable upsampling operation that again constructs high-resolution features and finally class labels based on the more abstract low-resolution representations learned at intermediate layers



Recap: Convolution Layers

- **Transposed convolution layers:** learnable upsampling operation
- **Recap:** normal convolution operation
 - Learnable $k \times k$ kernel matrix
 - Move kernel matrix across input to generate single element in output
 - Multiply kernel elements with input elements and sum up
 - Stride controls movement of kernel in input, stride > 1 reduces output size

Standard convolution: 6 x 6 input, 3 x 3 kernel, stride one

Input						Kernel			+0	Output			
0	0	0	0	0	0	1	0	-1		-2			
0	-1	2	2	1	0	2	1	0					
0	1	-1	1	2	0	0	1	2					
0	-1	0	2	1	0								
0	1	1	0	2	0								
0	0	0	0	0	0								

Recap: Convolution Layers

- **Transposed convolution layers:** learnable upsampling operation
- **Recap:** normal convolution operation
 - Learnable $k \times k$ kernel matrix
 - Move kernel matrix across input to generate single element in output
 - Multiply kernel elements with input elements and sum up
 - Stride controls movement of kernel in input, stride > 1 reduces output size

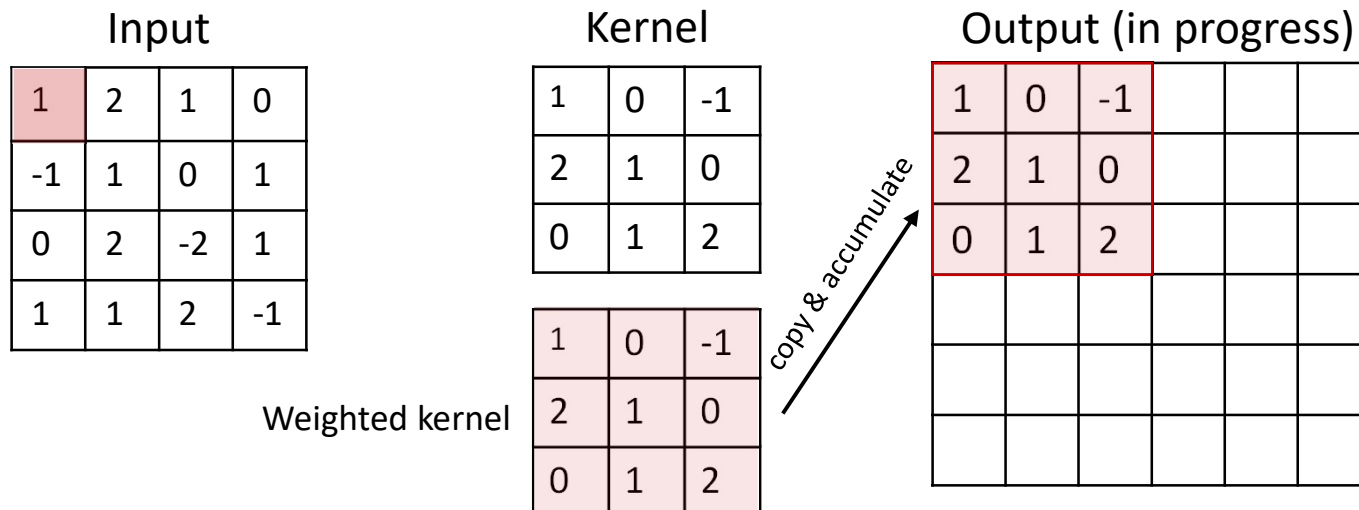
Standard convolution: 6 x 6 input, 3 x 3 kernel, stride one

Input						Kernel			Output			
0	0	0	0	0	0	1	0	-1	-2	1		
0	-1	2	2	1	0	2	1	0				
0	1	-1	1	2	0	0	1	2				
0	-1	0	2	1	0	+0						
0	1	1	0	2	0							
0	0	0	0	0	0							

Transposed Convolution Operation

- **Transposed convolution layers:** learnable upsampling operation
 - Learnable $k \times k$ kernel matrix
 - Move across single elements in input, multiply kernel matrix with input element and copy the "weighted" kernel matrix to the output
 - Overlapping outputs accumulate
 - Stride controls movement in output space, stride > 1 increases output size

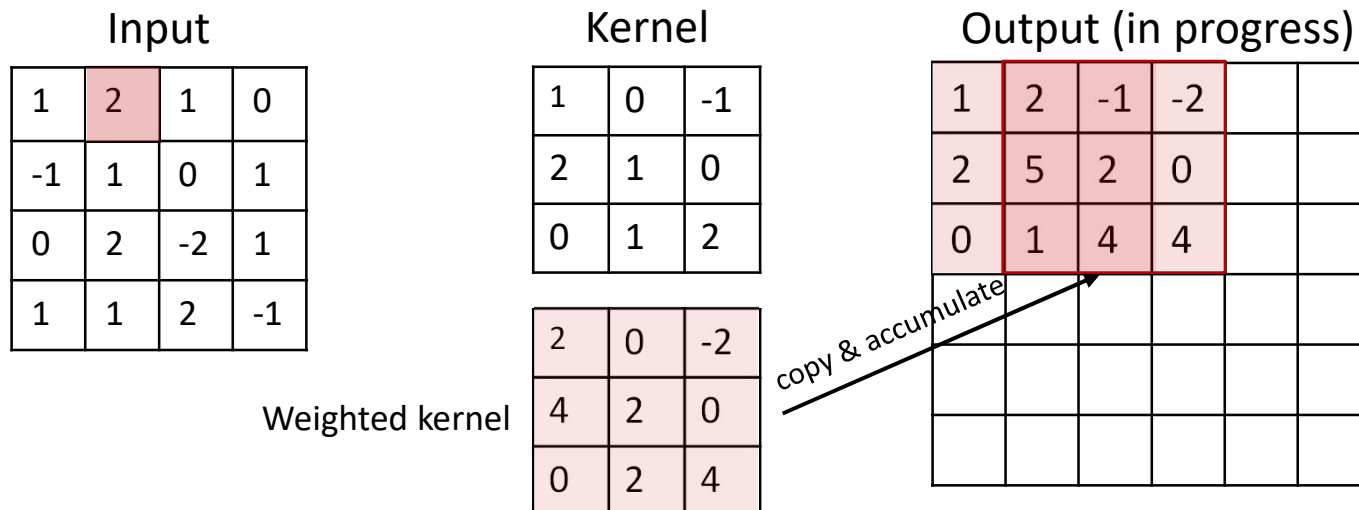
Transposed convolution: 4 x 4 input, 3 x 3 kernel, stride one



Transposed Convolution Operation

- **Transposed convolution layers:** learnable upsampling operation
 - Learnable $k \times k$ kernel matrix
 - Move across single elements in input, multiply kernel matrix with input element and copy the "weighted" kernel matrix to the output
 - Overlapping outputs accumulate
 - Stride controls movement in output space, stride > 1 increases output size

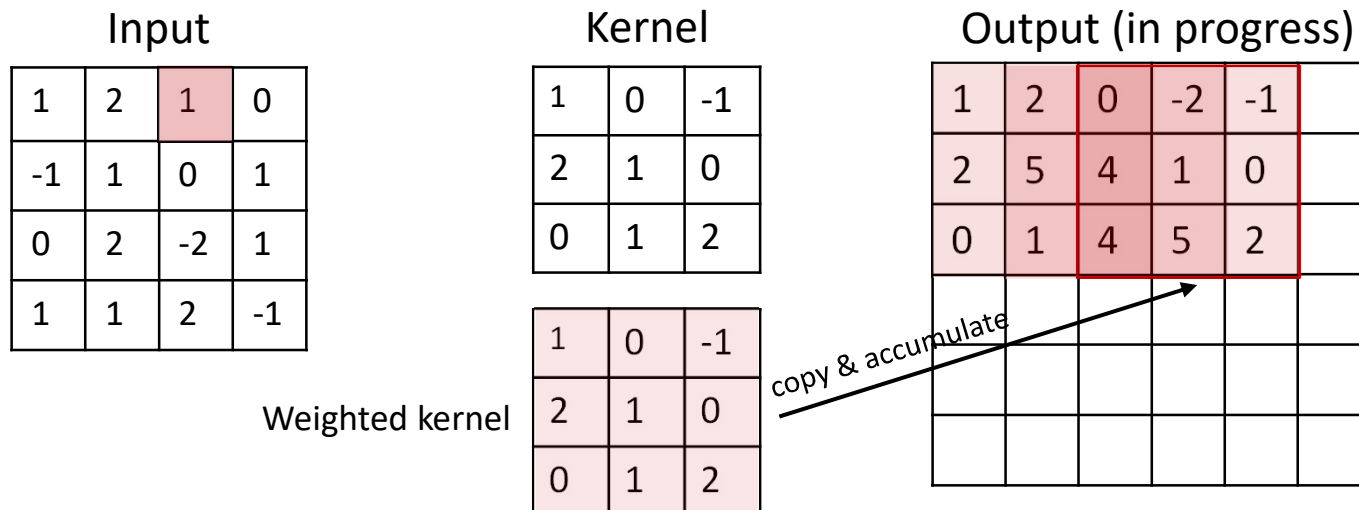
Transposed convolution: 4 x 4 input, 3 x 3 kernel, stride one



Transposed Convolution Operation

- **Transposed convolution layers:** learnable upsampling operation
 - Learnable $k \times k$ kernel matrix
 - Move across single elements in input, multiply kernel matrix with input element and copy the "weighted" kernel matrix to the output
 - Overlapping outputs accumulate
 - Stride controls movement in output space, stride > 1 increases output size


Transposed convolution: 4 x 4 input, 3 x 3 kernel, stride one



Transposed Convolution Operation

- **Transposed convolution layers:** learnable upsampling operation
 - Learnable $k \times k$ kernel matrix
 - Move across single elements in input, multiply kernel matrix with input element and copy the "weighted" kernel matrix to the output
 - Overlapping outputs accumulate
 - Stride controls movement in output space, stride > 1 increases output size


Transposed convolution: 4 x 4 input, 3 x 3 kernel, stride one



Transposed Convolution Operation

- **Transposed convolution layers:** learnable upsampling operation
 - Learnable $k \times k$ kernel matrix
 - Move across single elements in input, multiply kernel matrix with input element and copy the "weighted" kernel matrix to the output
 - Overlapping outputs accumulate
 - Stride controls movement in output space, stride > 1 increases output size


Transposed convolution: 4 x 4 input, 3 x 3 kernel, stride one



Transposed Convolution Operation

- **Transposed convolution layers:** learnable upsampling operation
 - Learnable $k \times k$ kernel matrix
 - Move across single elements in input, multiply kernel matrix with input element and copy the "weighted" kernel matrix to the output
 - Overlapping outputs accumulate
 - Stride controls movement in output space, stride > 1 increases output size

Transposed convolution: 4 x 4 input, 3 x 3 kernel, stride one



Transposed Convolution Operation: Padding

- **Padding for transposed convolution operation:** preserve size at stride=1
- For standard convolution, we pad the input with zeros to preserve size

Padded Input						Kernel			Output			
0	0	0	0	0	0	1	0	-1	-2			
0	-1	2	2	1	0	2	1	0				
0	1	-1	1	2	0	0	1	2				
0	-1	0	2	1	0							
0	1	1	0	2	0							
0	0	0	0	0	0							

+0

- For transposed convolutions, we crop the output (also called padding)

Input				Kernel			Output (in progress)					
1	2	1	0	1	0	-1	1	2	0	-2	-1	0
-1	1	0	1	2	1	0	1	6	5	0	0	0
0	2	-2	1	0	1	2	-2	2	5	5	2	0
1	1	2	-1				0	-1	-1	2		

Transposed Convolution Operation: Striding

- **Striding for transposed convolution operation:** actual upsampling
 - Striding s moves kernel in output space by s elements each time
 - If $s > 1$, this increases output size and performs upsampling

Example: Stride 2

Input				Kernel			Output (in progress)							
1	2	1	0	1	0	-1	1	0	-1					
-1	1	0	1	2	1	0	2	1	0					
0	2	-2	1	0	1	2	0	1	2					
1	1	2	-1											

Weighted kernel

1	0	-1
2	1	0
0	1	2

Transposed Convolution Operation: Striding

- **Striding for transposed convolution operation:** actual upsampling
 - Striding s moves kernel in output space by s elements each time
 - If $s > 1$, this increases output size and performs upsampling

Example: Stride 2

Input				Kernel			Output (in progress)							
1	2	1	0	1	0	-1	1	0	1	0	-2			
-1	1	0	1	2	1	0	2	1	4	2	0			
0	2	-2	1	0	1	2	0	1	2	2	4			
1	1	2	-1											
				2	0	-2								
				4	2	0								
				0	2	4								

Weighted kernel

Transposed Convolution Operation: Striding

- **Striding for transposed convolution operation:** actual upsampling
 - Striding s moves kernel in output space by s elements each time
 - If $s > 1$, this increases output size and performs upsampling

Example: Stride 2

Input				Kernel			Output (in progress)							
1	2	1	0	1	0	-1	1	0	1	0	-1	0	-1	
-1	1	0	1	2	1	0	2	1	4	2	2	1	0	
0	2	-2	1	0	1	2	0	1	2	2	4	1	2	
1	1	2	-1											

Weighted kernel

1	0	-1
2	1	0
0	1	2

Transposed Convolution Operation: Striding

- **Striding for transposed convolution operation:** actual upsampling
 - Striding s moves kernel in output space by s elements each time
 - If $s > 1$, this increases output size and performs upsampling

Example: Stride 2

Input				Kernel			Output (in progress)								
1	2	1	0	1	0	-1	1	0	1	0	-1	0	-1	0	0
-1	1	0	1	2	1	0	2	1	4	2	2	1	0	0	0
0	2	-2	1	0	1	2	0	1	2	2	4	1	2	0	0
1	1	2	-1												

Weighted kernel

0	0	0
0	0	0
0	0	0

Weighted kernel

Transposed Convolution Operation: Striding

- **Striding for transposed convolution operation:** actual upsampling
 - Striding s moves kernel in output space by s elements each time
 - If $s > 1$, this increases output size and performs upsampling

Example: Stride 2

Input				Kernel			Output (in progress)								
1	2	1	0	1	0	-1	1	0	1	0	-1	0	-1	0	0
-1	1	0	1	2	1	0	2	1	4	2	2	1	0	0	0
0	2	-2	1	0	1	2	-1	1	3	2	4	1	2	0	0
1	1	2	-1				-2	-1	0						
				-1	0	1	0	-1	-2						
				-2	-1	0									
				0	-1	-2									

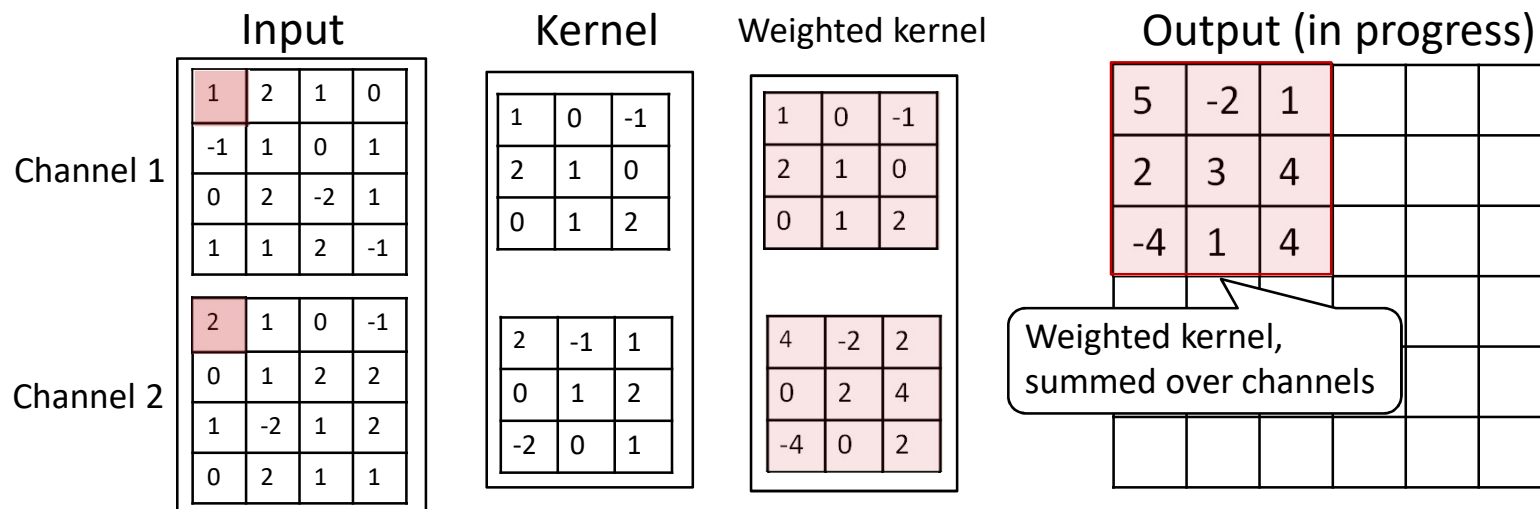
Weighted kernel

Weighted kernel

Transposed Convolutions: Input Channels

- As for normal convolution layers, we usually have multiple input channels for a transposed convolution layer
- As for normal convolution layers, in this case there is a further dimension in the kernel that corresponds to the different input channels
- Each channel in the kernel is multiplied with the element from corresponding channel in input. Results in output are accumulated over the channels

Example: Two input channels, stride 1



Transposed Convolutions: Input Channels


- As for normal convolution layers, we usually have multiple input channels for a transposed convolution layer
- As for normal convolution layers, in this case there is a further dimension in the kernel that corresponds to the different input channels
- Each channel in the kernel is multiplied with the element from corresponding channel in input. Results in output are accumulated over the channels

Example: Two input channels, stride 1

	Input	Kernel	Weighted kernel	Output (in progress)																																																																						
Channel 1	<table><tr><td>1</td><td>2</td><td>1</td><td>0</td></tr><tr><td>-1</td><td>1</td><td>0</td><td>1</td></tr><tr><td>0</td><td>2</td><td>-2</td><td>1</td></tr><tr><td>1</td><td>1</td><td>2</td><td>-1</td></tr></table>	1	2	1	0	-1	1	0	1	0	2	-2	1	1	1	2	-1	<table><tr><td>1</td><td>0</td><td>-1</td></tr><tr><td>2</td><td>1</td><td>0</td></tr><tr><td>0</td><td>1</td><td>2</td></tr></table>	1	0	-1	2	1	0	0	1	2	<table><tr><td>2</td><td>0</td><td>-2</td></tr><tr><td>4</td><td>2</td><td>0</td></tr><tr><td>0</td><td>2</td><td>4</td></tr></table>	2	0	-2	4	2	0	0	2	4	<table><tr><td>5</td><td>2</td><td>0</td><td>-1</td><td></td><td></td></tr><tr><td>2</td><td>7</td><td>7</td><td>2</td><td></td><td></td></tr><tr><td>-4</td><td>-1</td><td>6</td><td>5</td><td></td><td></td></tr><tr><td></td><td></td><td></td><td></td><td></td><td></td></tr><tr><td></td><td></td><td></td><td></td><td></td><td></td></tr><tr><td></td><td></td><td></td><td></td><td></td><td></td></tr></table>	5	2	0	-1			2	7	7	2			-4	-1	6	5																				
1	2	1	0																																																																							
-1	1	0	1																																																																							
0	2	-2	1																																																																							
1	1	2	-1																																																																							
1	0	-1																																																																								
2	1	0																																																																								
0	1	2																																																																								
2	0	-2																																																																								
4	2	0																																																																								
0	2	4																																																																								
5	2	0	-1																																																																							
2	7	7	2																																																																							
-4	-1	6	5																																																																							
Channel 2	<table><tr><td>2</td><td>1</td><td>0</td><td>-1</td></tr><tr><td>0</td><td>1</td><td>2</td><td>2</td></tr><tr><td>1</td><td>-2</td><td>1</td><td>2</td></tr><tr><td>0</td><td>2</td><td>1</td><td>1</td></tr></table>	2	1	0	-1	0	1	2	2	1	-2	1	2	0	2	1	1	<table><tr><td>2</td><td>-1</td><td>1</td></tr><tr><td>0</td><td>1</td><td>2</td></tr><tr><td>-2</td><td>0</td><td>1</td></tr></table>	2	-1	1	0	1	2	-2	0	1	<table><tr><td>2</td><td>-1</td><td>1</td></tr><tr><td>0</td><td>1</td><td>2</td></tr><tr><td>-2</td><td>0</td><td>1</td></tr></table>	2	-1	1	0	1	2	-2	0	1																																					
2	1	0	-1																																																																							
0	1	2	2																																																																							
1	-2	1	2																																																																							
0	2	1	1																																																																							
2	-1	1																																																																								
0	1	2																																																																								
-2	0	1																																																																								
2	-1	1																																																																								
0	1	2																																																																								
-2	0	1																																																																								


Transposed Convolutions: Output Channels

- As for normal convolution layers, we usually have multiple output channels for a transposed convolution layer
- As for normal convolution layers, each output channel is given by applying a different kernel to the same input in the same way. Resulting output channels are then stacked as usual.



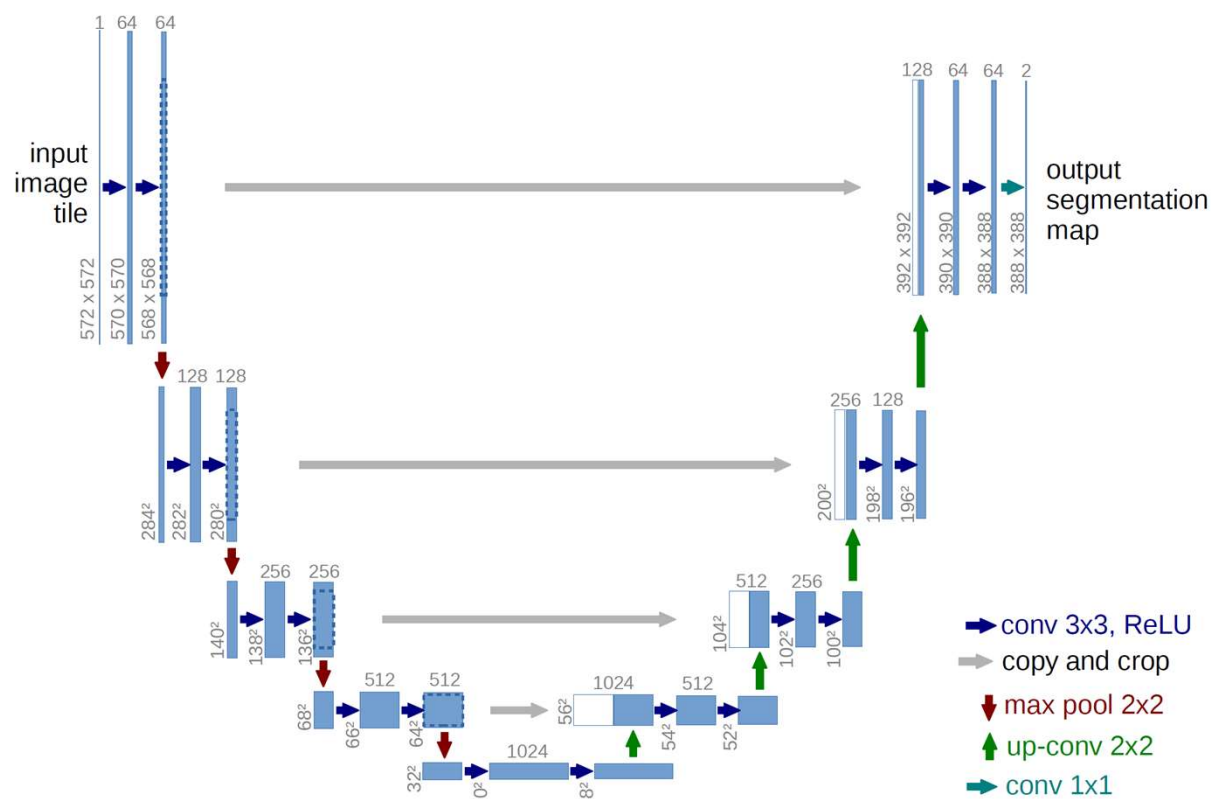
Problem Setting Segmentation

- **Back to full architecture:** Transposed convolution can be used for spatial upsampling
- Also called „deconvolution“, „upconvolution“, „fractionally strided convolution“
- Yields hourglass-shaped architecture with downsampling for learning about large-scale spatial context followed by upsampling for higher spatial resolution features learned based on the intermediate layers
- Upsampling step has learnable parameters, learns spatially high-resolution features that are helpful for the final segmentation



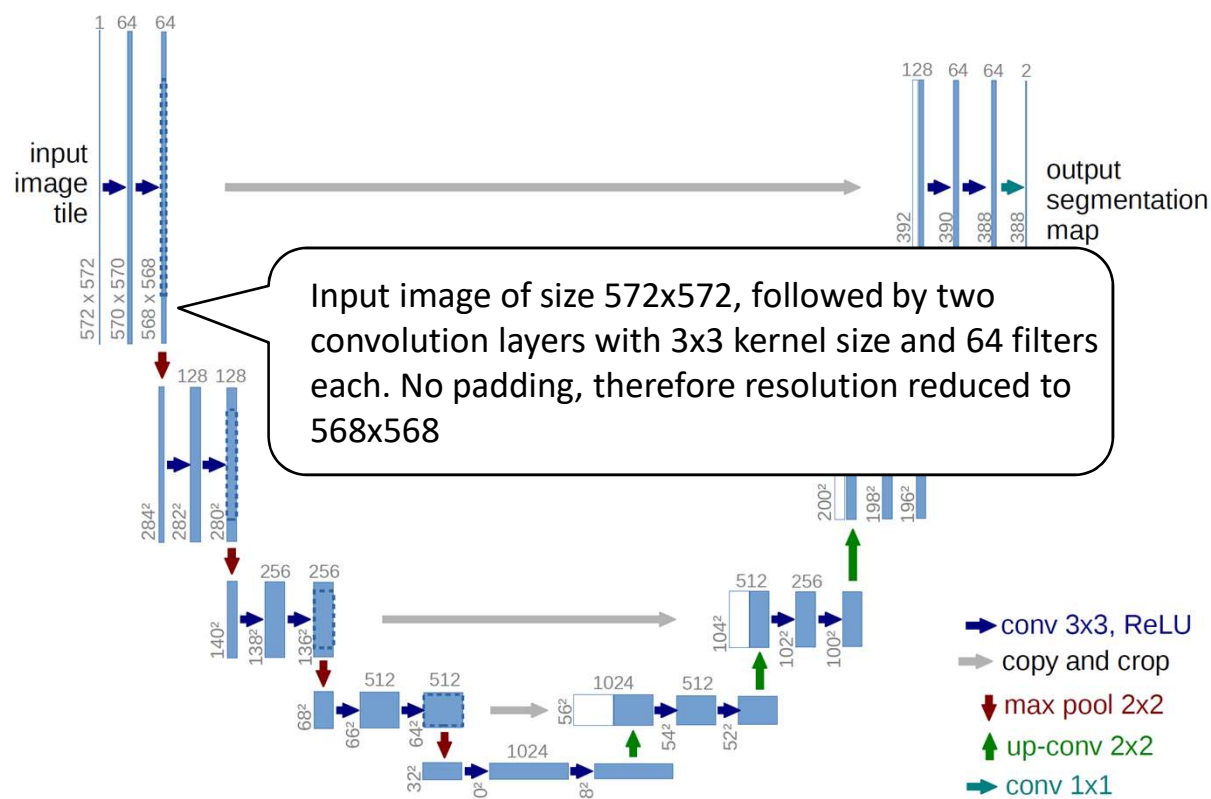
The U-Net Architecture for Segmentation

- U-Net architecture for segmentation [Ronneberger et al., 2015]: downsampling, upsampling, and direct connections between lower and higher layers



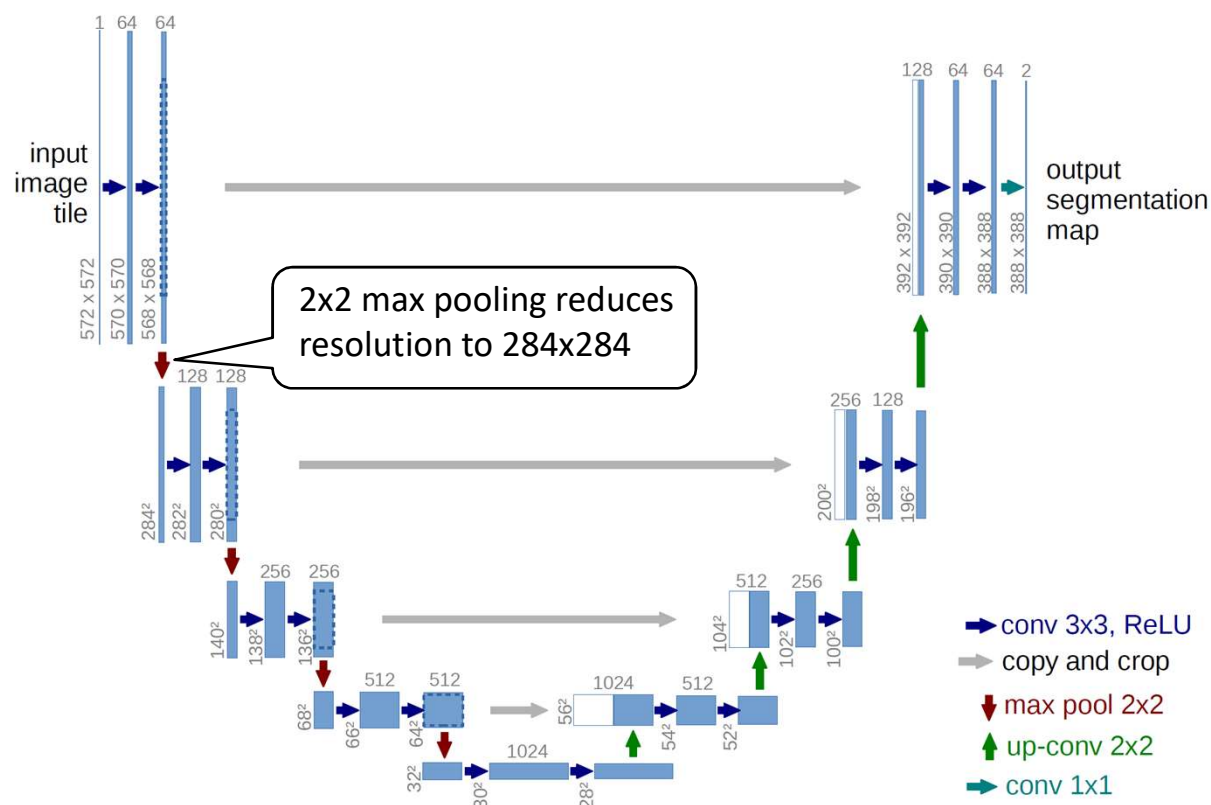
The U-Net Architecture for Segmentation

- U-Net architecture for segmentation [Ronneberger et al., 2015]: downsampling, upsampling, and direct connections between lower and higher layers




The U-Net Architecture for Segmentation

- U-Net architecture for segmentation [Ronneberger et al., 2015]: downsampling, upsampling, and direct connections between lower and higher layers




The U-Net Architecture for Segmentation

- U-Net architecture for segmentation [Ronneberger et al., 2015]: downsampling, upsampling, and direct connections between lower and higher layers




The U-Net Architecture for Segmentation

- U-Net architecture for segmentation [Ronneberger et al., 2015]: downsampling, upsampling, and direct connections between lower and higher layers




The U-Net Architecture for Segmentation

- U-Net architecture for segmentation [Ronneberger et al., 2015]: downsampling, upsampling, and direct connections between lower and higher layers



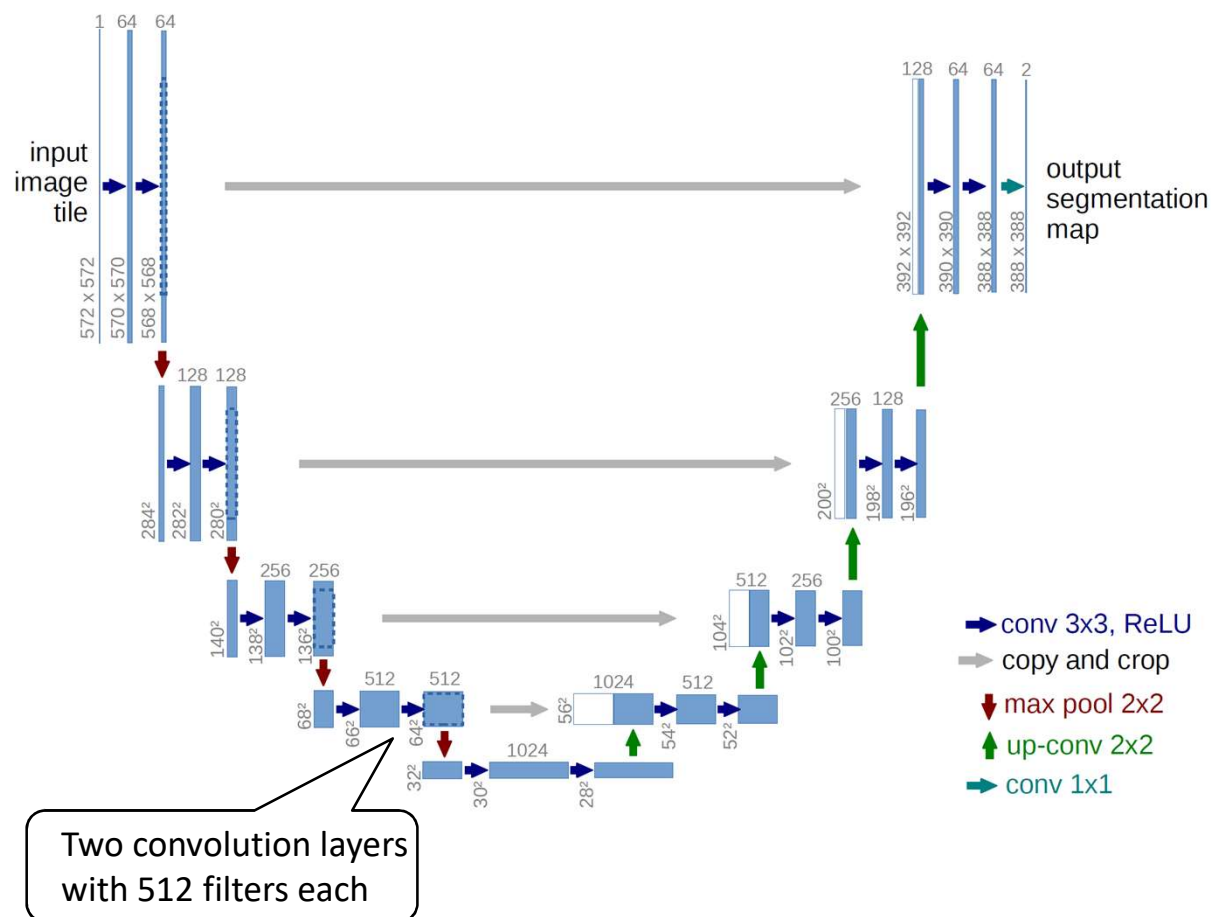
The U-Net Architecture for Segmentation

- U-Net architecture for segmentation [Ronneberger et al., 2015]: downsampling, upsampling, and direct connections between lower and higher layers



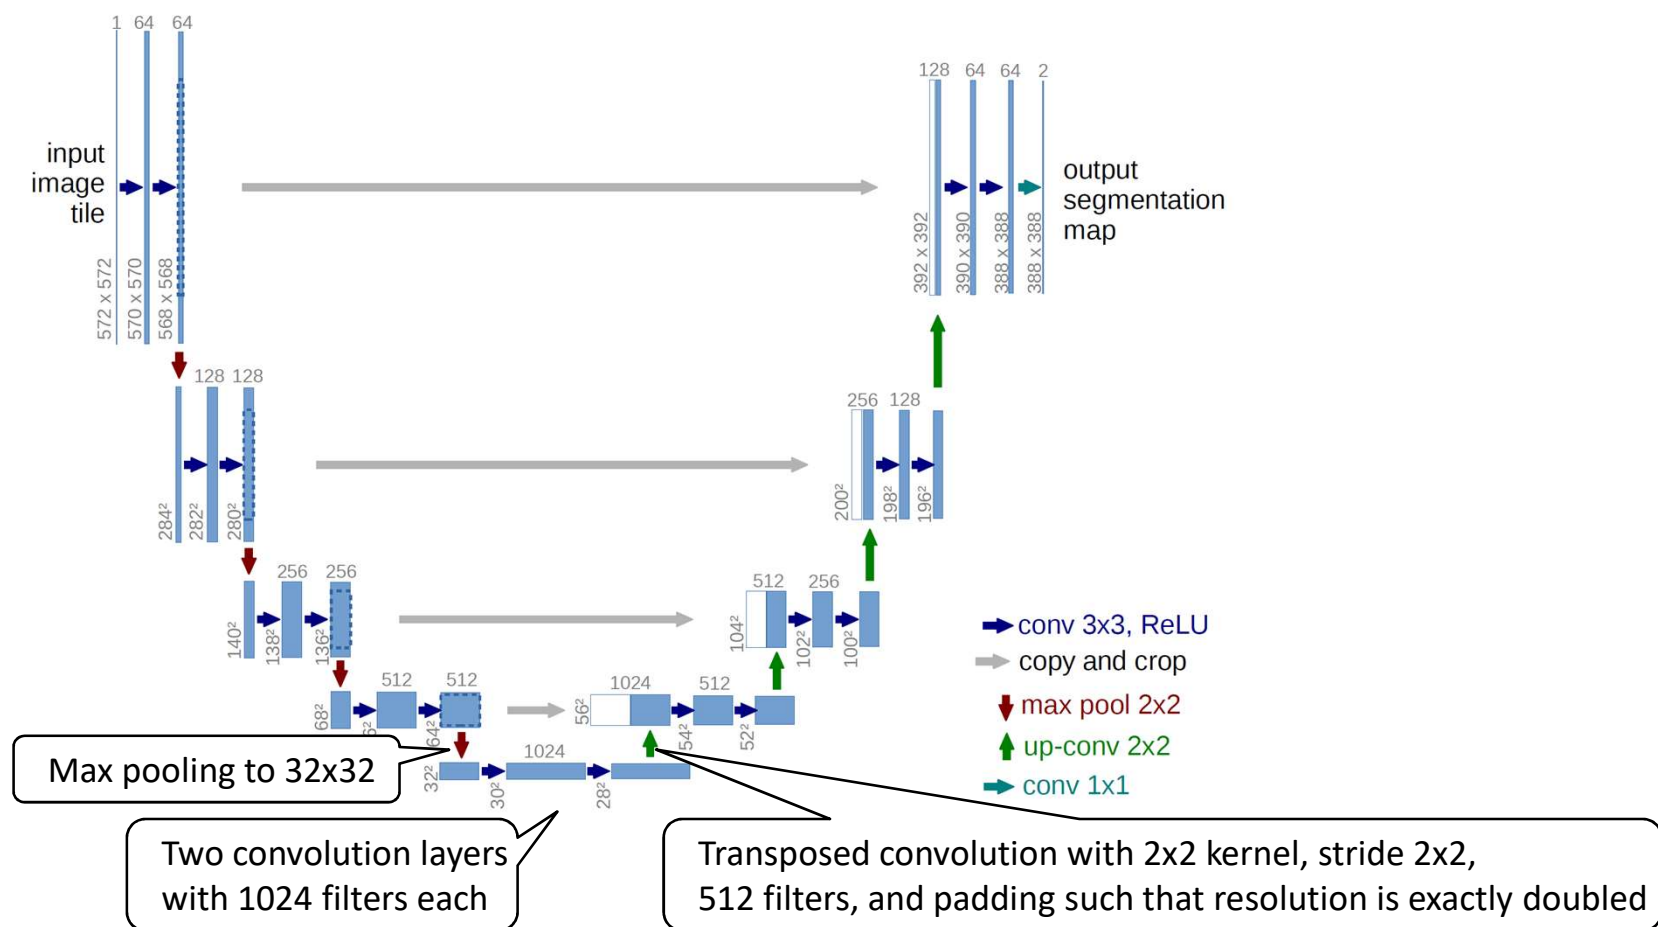
The U-Net Architecture for Segmentation

- U-Net architecture for segmentation [Ronneberger et al., 2015]: downsampling, upsampling, and direct connections between lower and higher layers



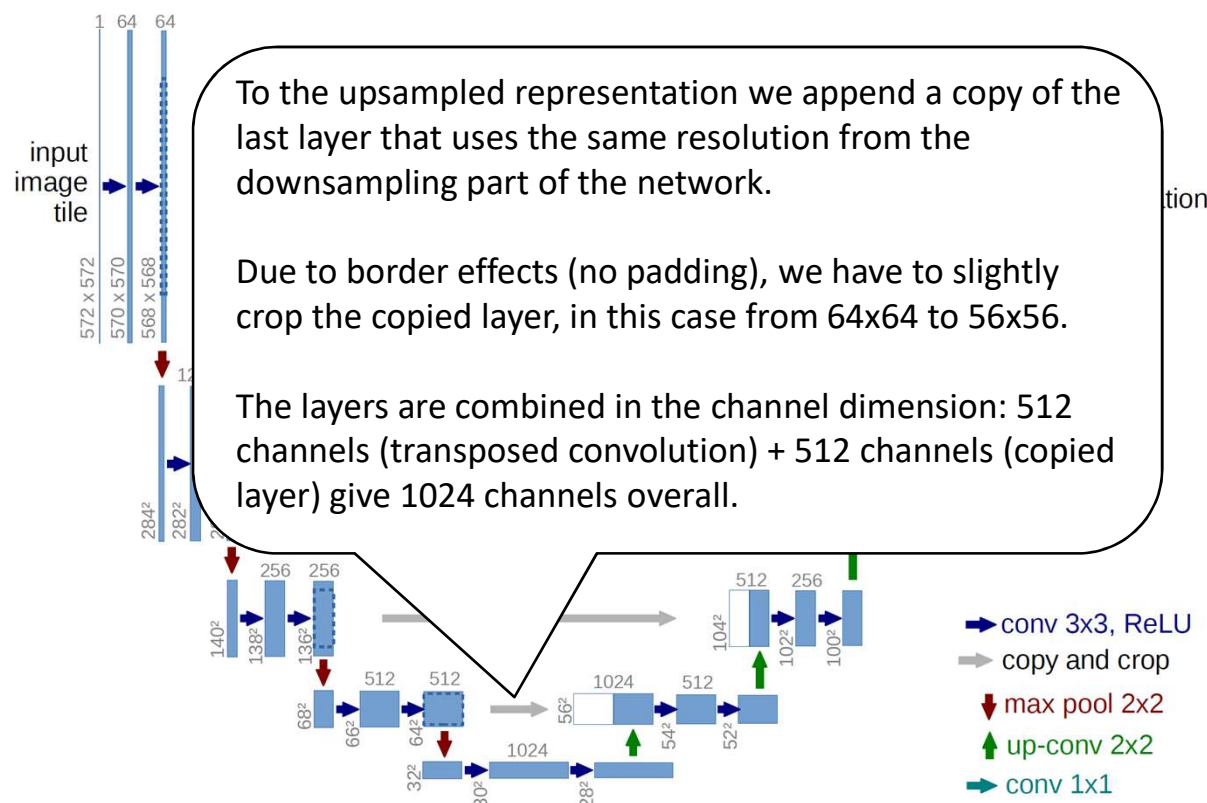
The U-Net Architecture for Segmentation

- U-Net architecture for segmentation [Ronneberger et al., 2015]: downsampling, upsampling, and direct connections between lower and higher layers



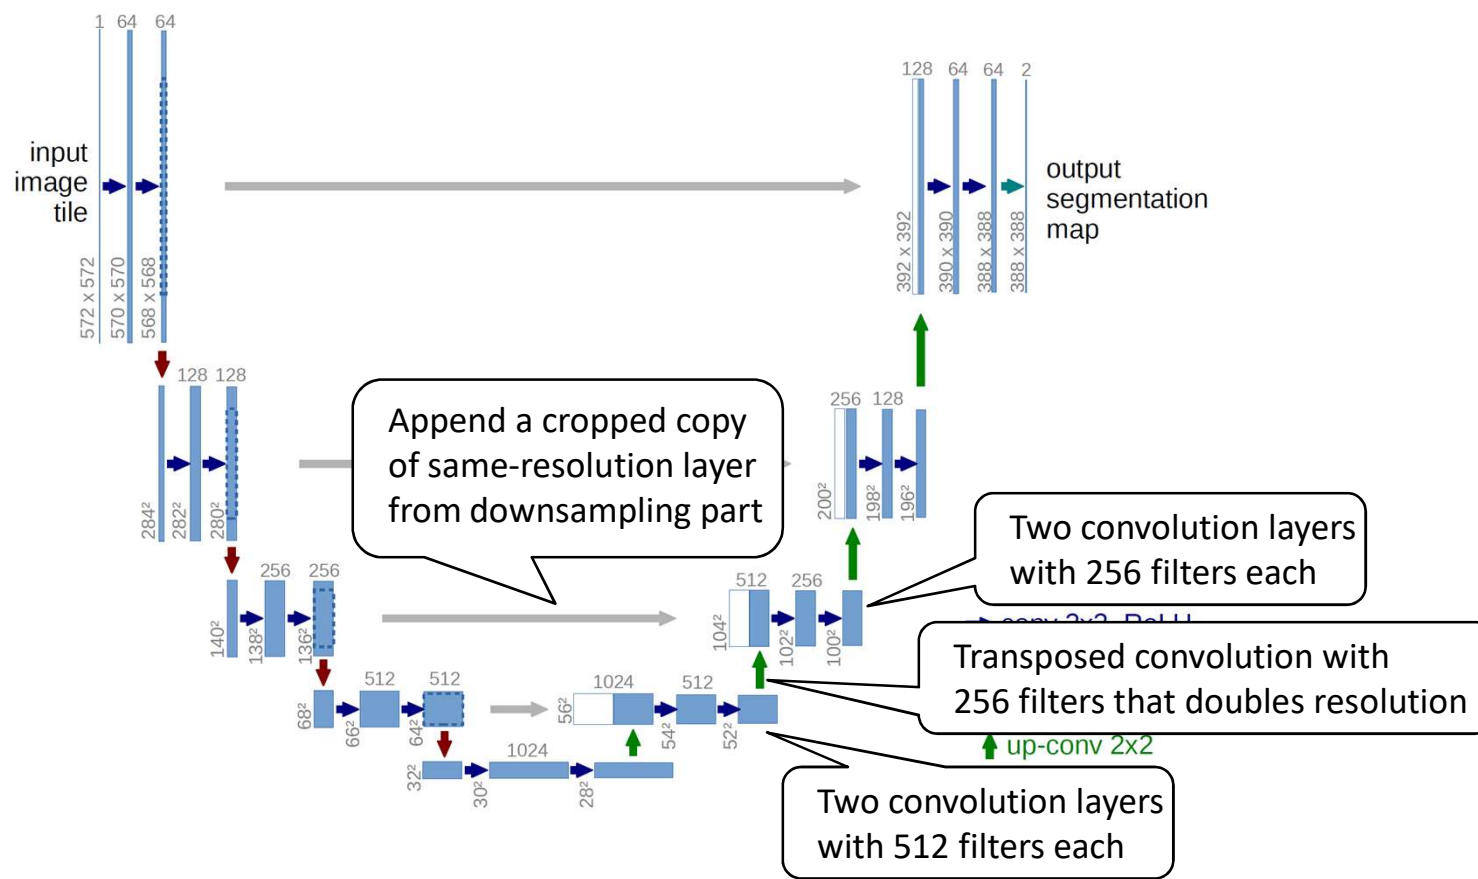
The U-Net Architecture for Segmentation

- U-Net architecture for segmentation [Ronneberger et al., 2015]: downsampling, upsampling, and direct connections between lower and higher layers



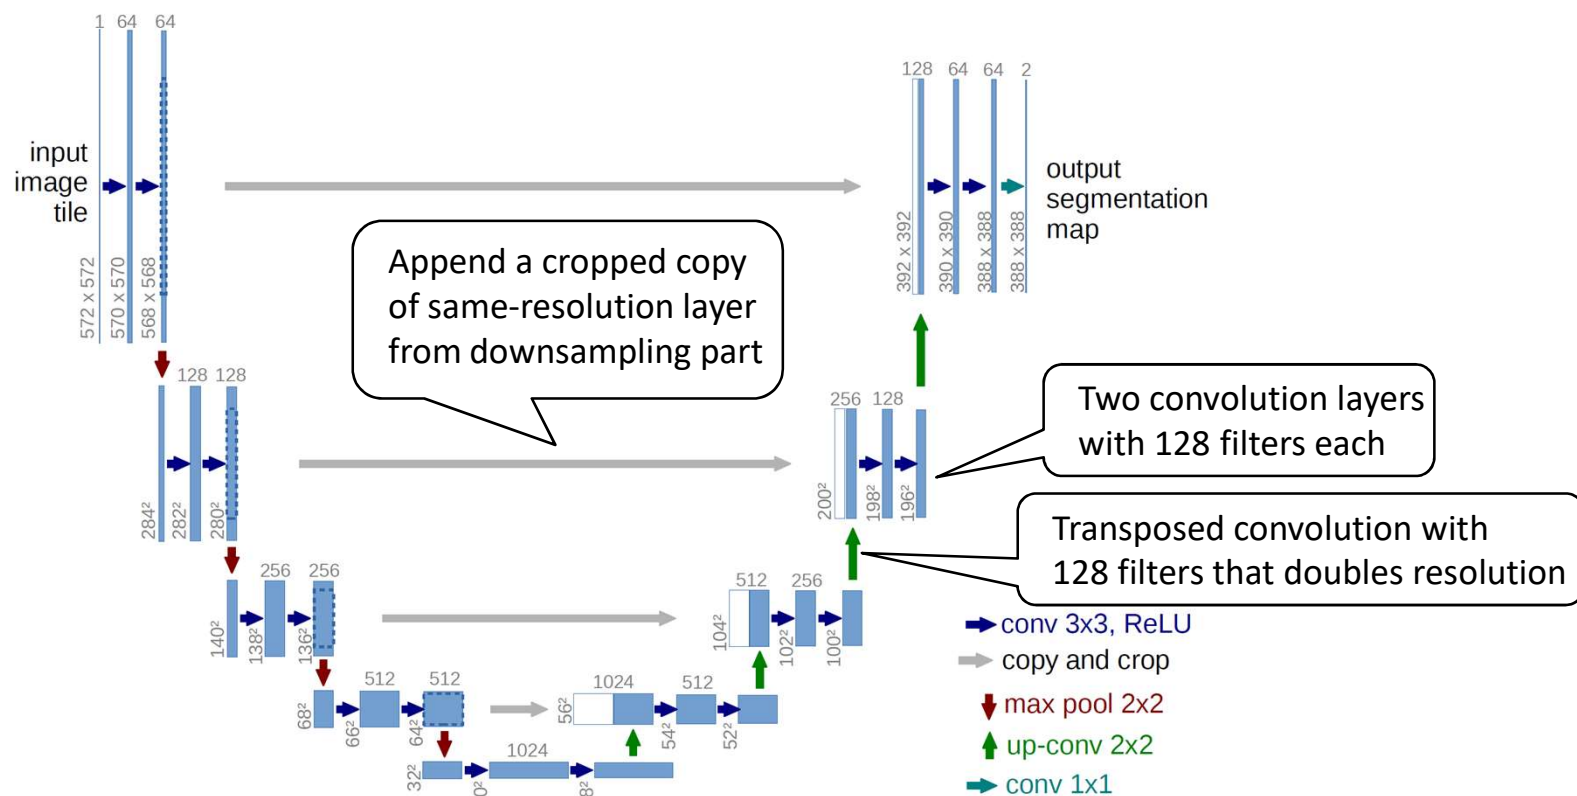
The U-Net Architecture for Segmentation

- U-Net architecture for segmentation [Ronneberger et al., 2015]: downsampling, upsampling, and direct connections between lower and higher layers



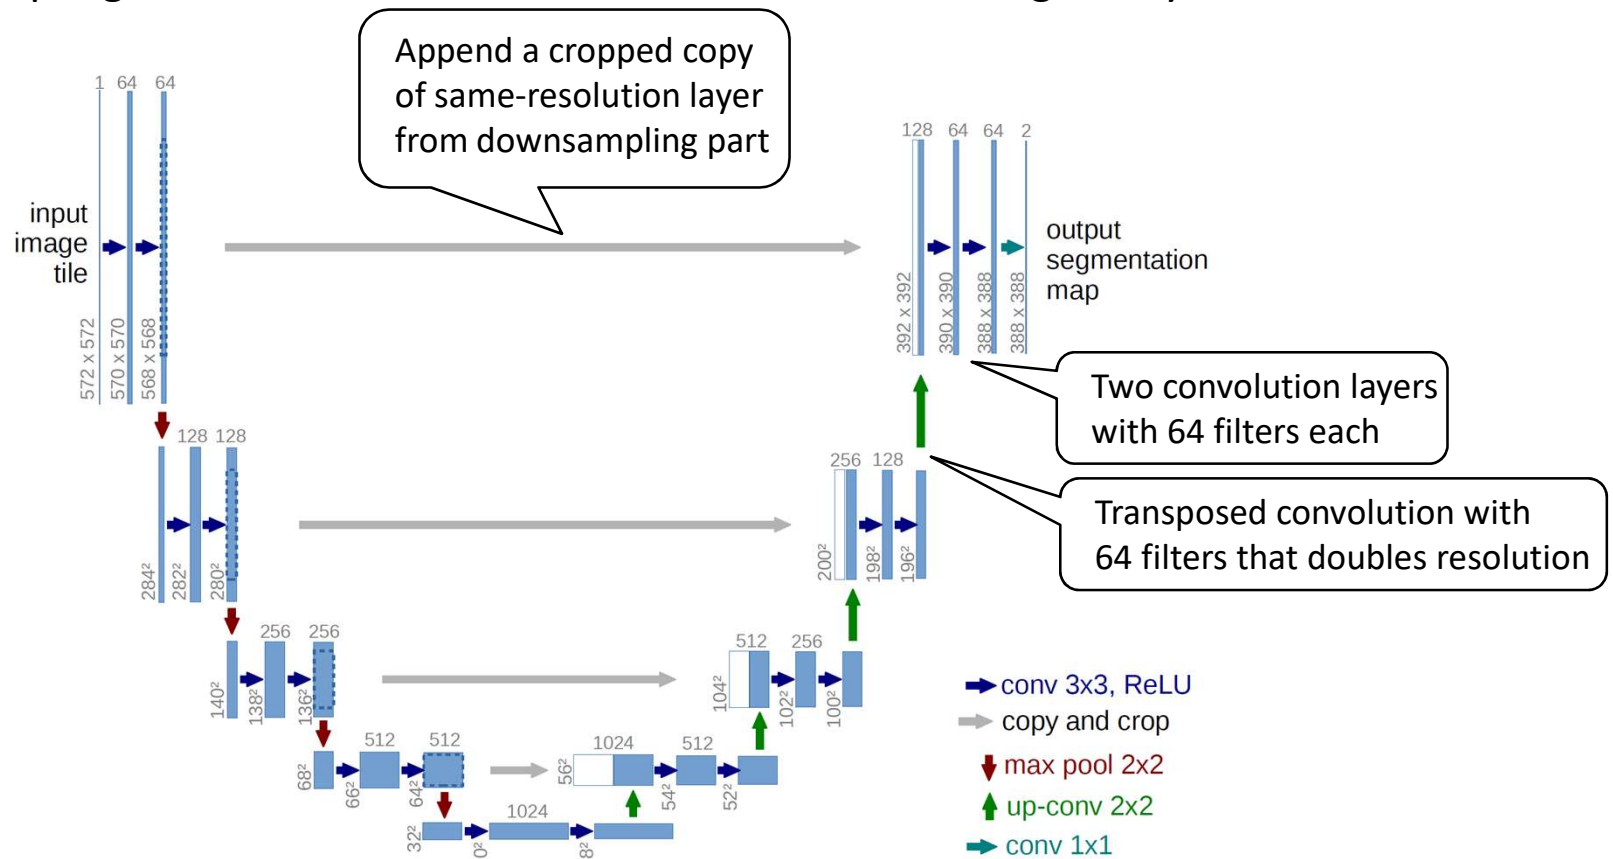
The U-Net Architecture for Segmentation

- U-Net architecture for segmentation [Ronneberger et al., 2015]: downsampling, upsampling, and direct connections between lower and higher layers



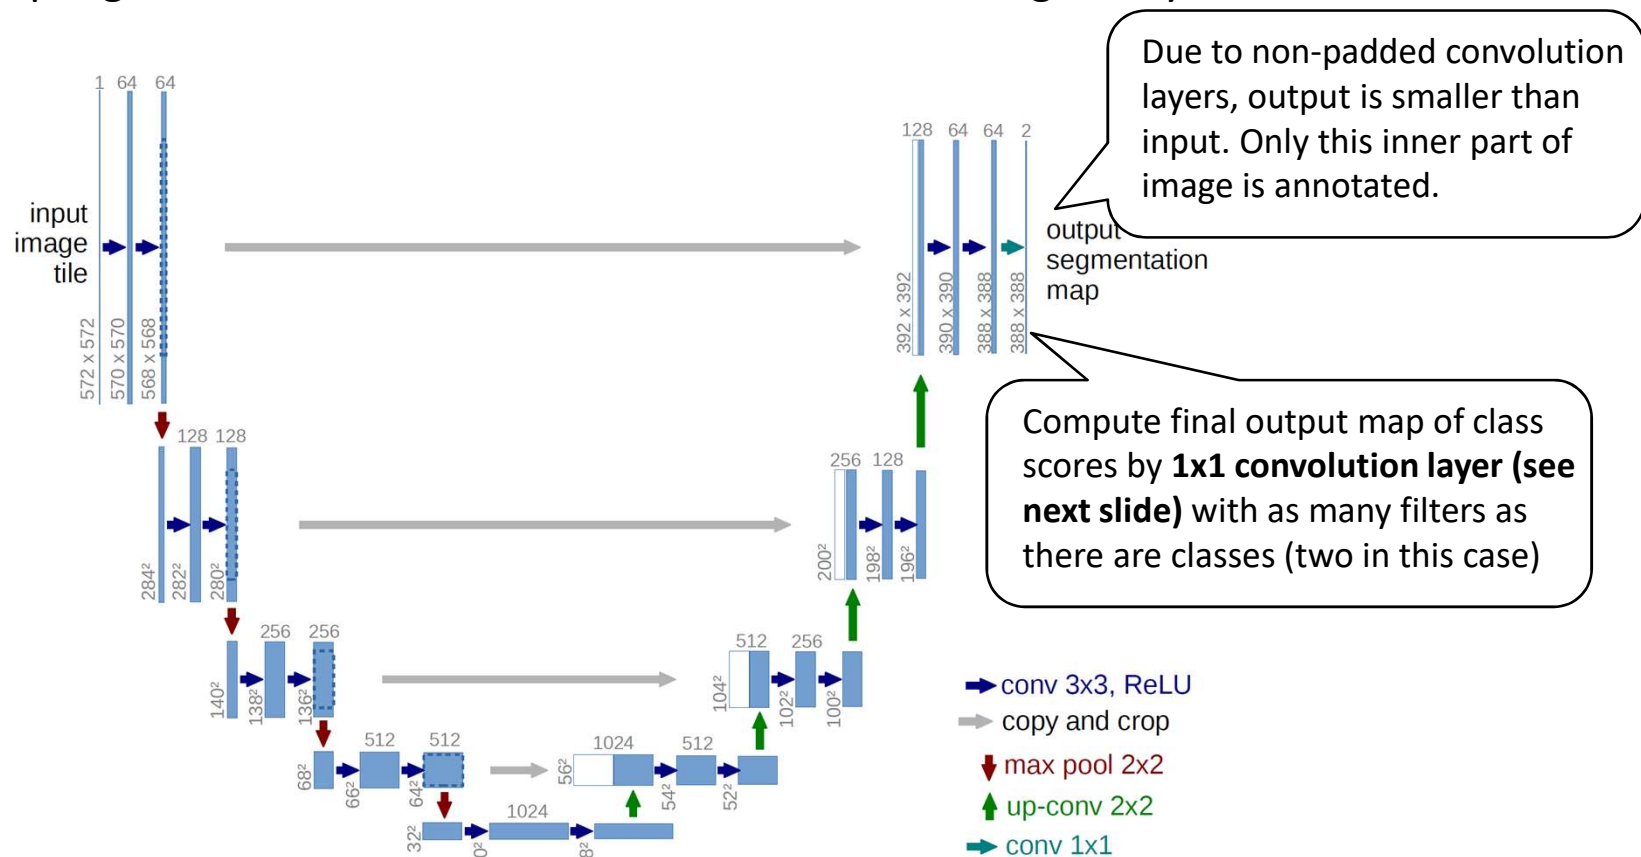
The U-Net Architecture for Segmentation

- U-Net architecture for segmentation [Ronneberger et al., 2015]: downsampling, upsampling, and direct connections between lower and higher layers



The U-Net Architecture for Segmentation

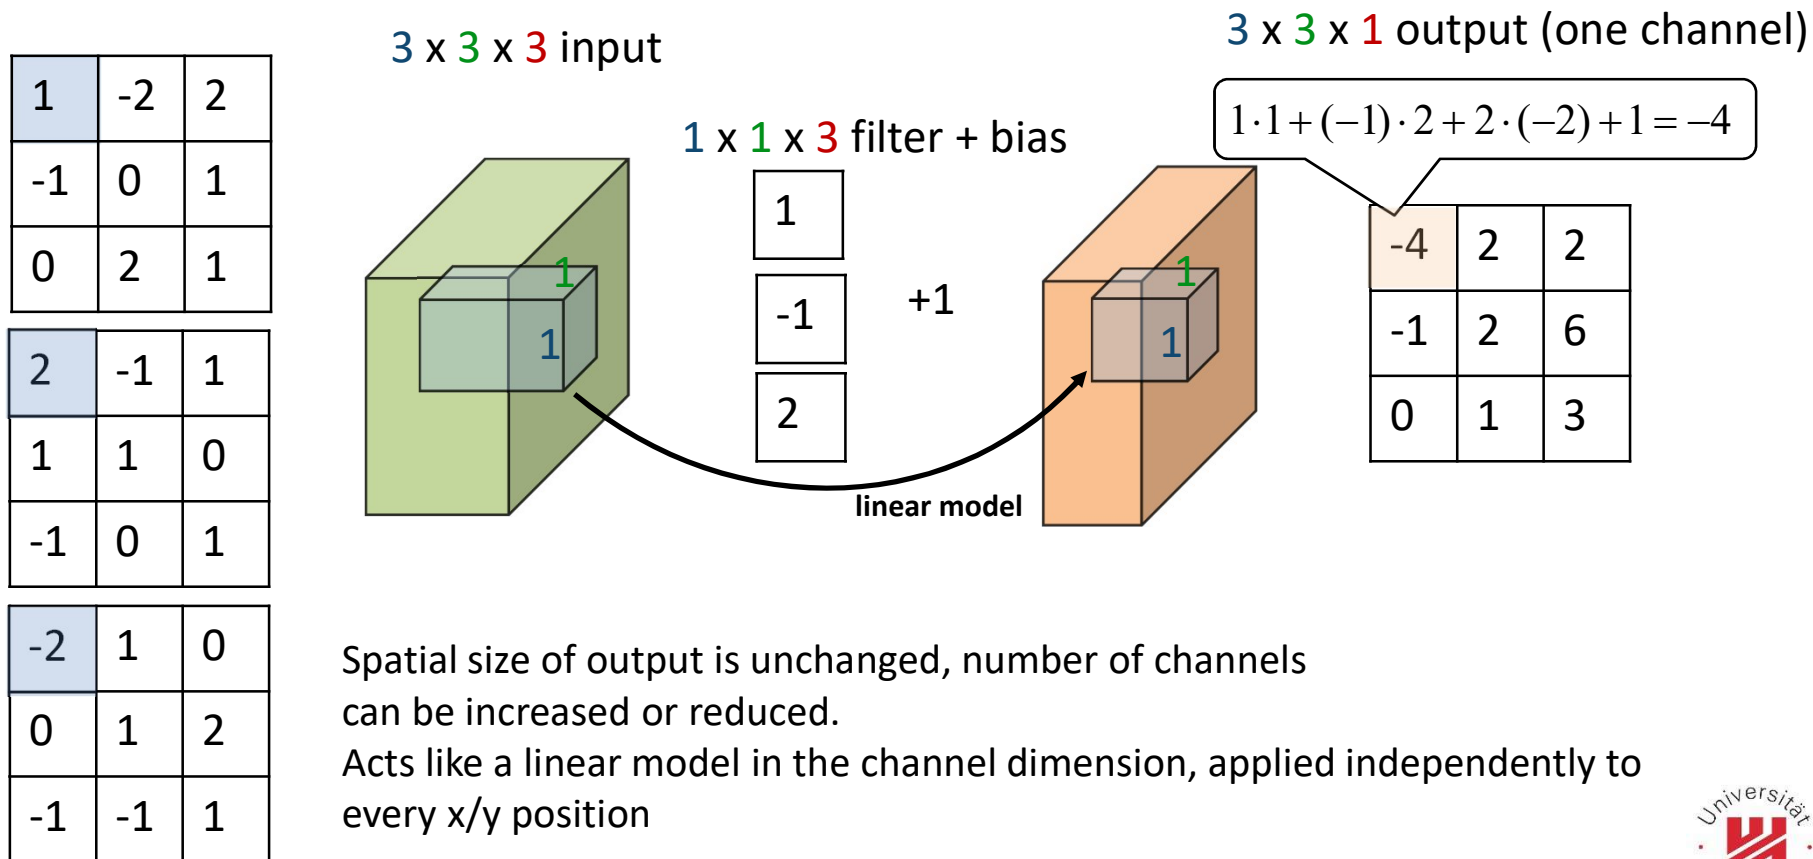
- U-Net architecture for segmentation [Ronneberger et al., 2015]: downsampling, upsampling, and direct connections between lower and higher layers



Special Case 1x1 Convolution

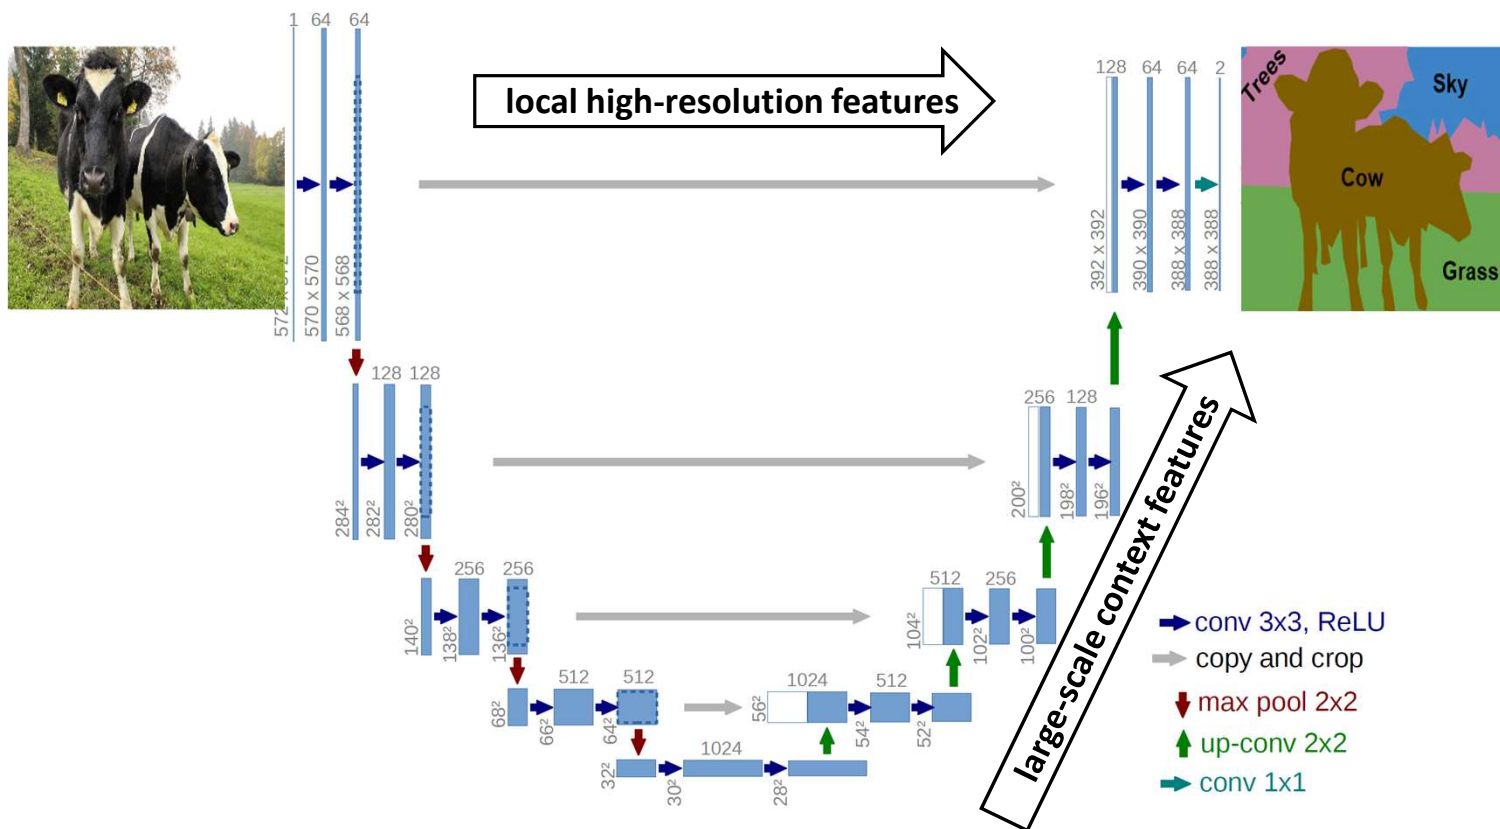
- Reminder: A frequently used special case of the standard convolution operation is a convolution with a 1x1 kernel size (and stride 1)

Example with three input channels



The U-Net Architecture for Segmentation

- The „copy“-pathes in UNet combine the information from the higher-resolution representations in the original image and the early layers with the more large-scale context from the intermediate layers



U-Net Loss Function for Segmentation

- Model is trained on annotated images, that is, pairs of images and class maps:
 $\mathbf{X} = \{\mathbf{x}_1, \dots, \mathbf{x}_n\}, \mathbf{Y} = \{\mathbf{y}_1, \dots, \mathbf{y}_n\}$ with $\mathbf{x}_i \in \mathbb{R}^{m \times l \times d}$, $\mathbf{y}_i \in \mathbb{R}^{m \times l}$
- Let $\mathbf{s}_i \in \mathbb{R}^{m \times l \times k}$ denote the final output tensor of the model (class scores for all pixels in output) for input \mathbf{x}_i .
- For each position x, y : compute a predicted class distribution by

$$p(\mathbf{y}[x, y] = j \mid \mathbf{x}_i, \boldsymbol{\theta}) = \frac{\exp(\mathbf{s}_i[x, y, j])}{\sum_{j'=1}^k \exp(\mathbf{s}_i[x, y, j'])}$$

Probability for class j according to model


All model parameters

- Loss for an image is cross entropy summed over positions in output


$$\ell(f_{\boldsymbol{\theta}}(\mathbf{x}_i), \mathbf{y}_i) = - \sum_{x, y} \log p(\mathbf{y}[x, y] = \mathbf{y}_i[x, y] \mid \mathbf{x}_i, \boldsymbol{\theta})$$

Summary: Segmentation, U-Net Architecture

- Segmentation problem:



- U-Net Architecture: downsampling for large-scale context, upsampling for final segmentation map, „copy“-connections to reuse local features



- As usual, gradient of entire model can be derived via automatic differentiation, and model can be trained end-to-end by stochastic gradient descent