

$$\begin{aligned}
 \text{Task 1)} \quad E[(\hat{R} - R)^2] &= E[(\hat{R} - E(\hat{R}) + E(\hat{R}) - R)^2] \\
 &= E[(\hat{R} - E(\hat{R}))^2 + (E(\hat{R}) - R)^2 + 2(\hat{R} - E(\hat{R}))(E(\hat{R}) - R)] \\
 &= E[(\hat{R} - E(\hat{R}))^2] + E[(E(\hat{R}) - R)^2] + 2E[(\hat{R} - E(\hat{R}))(E(\hat{R}) - R)] \\
 &= \text{Var}[\hat{R}] + \text{Bias}[\hat{R}]^2 + 2\left[E(\hat{R}^2) - E(\hat{R})E(R) - E(\hat{R})^2 + E(\hat{R})E(R)\right] \\
 &= \text{Var}[\hat{R}] + \text{Bias}[\hat{R}]^2
 \end{aligned}$$

$$\begin{aligned}
 \text{Task 3)} \\
 \hat{R}_T(f_{\theta^*}) &= \frac{1}{N} \sum_{n=1}^N l_{\text{eval}}(\bar{y}_n, f_{\theta^*}(\bar{x}_n)), \quad l_{\text{eval}} = \begin{cases} 0 & y = f_{\theta^*}(x) \\ 1 & y \neq f_{\theta^*}(x) \end{cases} \\
 \Rightarrow \hat{R}_T(f_{\theta^*}) &= \frac{1}{10} (0 \times 0 + 2 \times 4) = 0.4
 \end{aligned}$$

$$s_{R_T}^2 = \frac{\hat{R}_T(f_{\theta^*})(1 - \hat{R}_T(f_{\theta^*}))}{N} = \frac{0.4 \times 0.6}{10} = 0.024$$

for a two-banded confidence interval $1 - 2\delta = 0.95$

$$\Rightarrow \delta = 0.05 \rightarrow \Phi^{-1} = 1.96$$

$$\varepsilon = s_{R_T} \cdot \Phi^{-1}(1 - 2\delta) = \sqrt{0.024} \cdot (1.96) = 0.303$$

$$\Rightarrow \text{CI} = [\hat{R}_T(f_{\theta^*}) - \varepsilon, \hat{R}_T(f_{\theta^*}) + \varepsilon] =$$

$$[0.4 - 0.303, 0.4 + 0.303] = [0.097, 0.703]$$

Task 2

```
In [ ]: import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics import mean_squared_error

# Input: number of samples N
# Output: one-dimensional data set of N points where  $y_n = \sin(2 \pi x_n) + \epsilon_n$  as in lecture
def sine_data_set(N):
    np.random.seed(1234)
    x = np.random.uniform(0,1,(N))
    y = np.sin(x*2*np.pi)+np.random.normal(scale=0.2,size=(N))
    return x,y

# Input: one-dimensional inputs x as vector of length N, polynomial degree d
# Output: polynomial feature representation of the inputs as N x (d+1) matrix
def poly_features(x,d):
    X = np.zeros((x.shape[0],d+1))
    for i in range(0,d+1):
        X[:,i] = np.power(x,i)
    return X

# Input: instances X as N x M matrix, Labels y as vector of length N, Lambda for regularization
# Output: Learned parameter vector
def fit_ridge_regression(X,y,lambda_param):
    N = X.shape[0]
    M = X.shape[1]
    return np.linalg.solve(X.T @ X + N*lambda_param*np.eye(M), X.T @ y)

# Input: instances X as N x M matrix, model theta
# Output: predictions of model theta on X
def predict_regression(X,theta):
    return X @ theta
```

```

# Input: instances X, Labels y, number of cross-validation folds K, current fold k
# Output: train and test sets for fold k
def crossval_split(X,y,K,k):
    x_parts = np.array_split(X, K)
    y_parts = np.array_split(y, K)
    X_test = x_parts[k]
    y_test = y_parts[k]
    x_parts.pop(k)
    y_parts.pop(k)
    X_train = np.concatenate(x_parts)
    y_train = np.concatenate(y_parts)
    return X_train, y_train, X_test, y_test

```

```

In [ ]: N = 20
x,y = sine_data_set(N)
plt.xlim([0,1])
plt.ylim([-1.2,1.2])
plt.scatter(x,y)
def GridSearch_d(x,y):
    lamb = 0
    best_loss = float('inf')
    best_model_theta = None
    best_d = None
    loss_history = []
    K = 4
    for d in range(0, 11):
        losses = 0
        for k in range(K):
            X_train, y_train, X_test, y_test = crossval_split(x,y,K ,k)
            X = poly_features(X_train,d)
            X_TEST = poly_features(X_test,d)
            theta = fit_ridge_regression(X,y_train,lamb)
            y_predict = predict_regression(X_TEST,theta)
            mse = mean_squared_error(y_test, y_predict)
            losses += mse
        losses /= K
        loss_history.append(losses)
        if losses < best_loss:
            best_loss = losses
            best_model_theta = theta
            best_d = d
    print(f'best d is : {best_d}')

```


```


X = poly_features(x,best_d)
theta = fit_ridge_regression(X,y,lamb)
grid = np.arange(0,1,0.001)
plt.plot(grid,predict_regression(poly_features(grid,best_d),theta),'y')
fig, ax = plt.subplots()
ax.plot(range(0, 11), loss_history)
ax.set_ylim(0, 1)
plt.show()

```

GridSearch_d(x, y)

best d is : 3





```
In [ ]: N = 20
x,y = sine_data_set(N)
plt.xlim([0,1])
plt.ylim([-1.2,1.2])
plt.scatter(x,y)
def GridSearch_lamb(x,y):
    d = 10
    lambda_values = np.logspace(0, -9, num=10, base=10)
    best_loss = float('inf')
    best_model_theta = None
    best_lamb = None
    loss_history = []
    K = 4
    for lamb in lambda_values:
        losses = 0
        for k in range(K):
```


```


X_train, y_train, X_test, y_test = crossval_split(x,y,K ,k)
X = poly_features(X_train,d)
X_TEST = poly_features(X_test,d)
theta = fit_ridge_regression(X,y_train,lamb)
y_predict = predict_regression(X_TEST,theta)
mse = mean_squared_error(y_test, y_predict)
losses += mse
losses /= K
loss_history.append(losses)
if losses < best_loss:
    best_loss = losses
    best_model_theta = theta
    best_lamb = lamb
print(f'best lambda is : {best_lamb}')
X = poly_features(x,d)
theta = fit_ridge_regression(X,y,best_lamb)
grid = np.arange(0,1,0.001)
plt.plot(grid,predict_regression(poly_features(grid,10),theta),'y')
fig, ax = plt.subplots()
ax.plot(lambda_values, loss_history)
ax.set_ylim(0, 1)
plt.show()

```

GridSearch_lamb(x, y)

best lambda is : 1e-06





In []: