

Task I)

$$\textcircled{1} P(F=1|v) = \frac{\alpha_h + N_h}{\alpha_h + \alpha_t + N_h + N_t} = \frac{1+4}{1+1+4+3} = \frac{5}{9}$$

$$\textcircled{2} P(b=1|v) = \frac{\alpha_h + N_h}{\alpha_h + \alpha_t + N_h + N_t} = \frac{1+4}{1+1+4+3} = \frac{5}{9}$$

$$\textcircled{3} P(c=1|F=1, b=1, v) = \frac{\alpha_h + N_h}{\alpha_h + \alpha_T + N_h + N_t} = \frac{1+2}{1+1+2+0} = \frac{3}{4}$$

$$\textcircled{4} P(c=1|F=1, b=0, v) = \frac{\alpha_h + N_h}{\alpha_h + \alpha_T + N_h + N_t} = \frac{1+1}{1+1+1+1} = \frac{2}{4}$$

$$\textcircled{5} P(c=1|F=0, b=1, v) = \frac{\alpha_h + N_h}{\alpha_h + \alpha_T + N_h + N_t} = \frac{1+1}{1+1+1+1} = \frac{2}{4}$$

$$\textcircled{6} P(c=1|F=0, b=0, v) = \frac{\alpha_h + N_h}{\alpha_h + \alpha_T + N_h + N_t} = \frac{1+0}{1+1+0+1} = \frac{1}{3}$$

Task II)

@ Loss function for probabilistic regression is the likelihood which we try to maximize as:

$$\theta^* = \arg \max_{\theta} P(y|X, \theta) = \arg \max_{\theta} \prod_{n=1}^N P(y_n | x_n, \theta)$$

then we apply log: $\theta^* = \arg \min_{\theta} -\log P(y|X, \theta)$ (I)
which is strictly monotone and multiply by -1

$$\text{And in our case: } P(y|x, \theta) = N(y | x^T \theta, \sigma^2) = \frac{1}{\sqrt{2\pi} \sigma} e^{-\frac{(y - x^T \theta)^2}{2\sigma^2}} \quad \textcircled{II}$$

$$\begin{aligned}
 \textcircled{I} \text{ , } \textcircled{II} &\Rightarrow L(\theta) = \sum_{n=1}^N \log p(y_n | x_n, \theta) \\
 &= \sum_{n=1}^N \log \left[\frac{1}{\sigma^2 \sqrt{2\pi}} e^{-\frac{(y_n - x_n^T \theta)^2}{2\sigma^2}} \right] \\
 &= \sum_{n=1}^N \log \left(\frac{1}{\sigma^2 \sqrt{2\pi}} \right) + \sum_{n=1}^N \log \left[e^{-\frac{(y_n - x_n^T \theta)^2}{2\sigma^2}} \right] \\
 &= \underbrace{-N \log(\sigma^2 \sqrt{2\pi})}_{\text{constant}} - \underbrace{\frac{1}{2\sigma^2} \sum_{i=1}^N (y_i - x_i^T \theta)^2}_{\text{MSE}}
 \end{aligned}$$

Hence it is the same as minimizing the MSE.

$$\textcircled{b} \quad p(\theta) = \mathcal{N}(\theta | 0, \sigma^2 I) \quad \rightarrow \quad p(y | x, \theta) = \mathcal{N}(y | x^T \theta, \sigma^2)$$

$$p(\theta | y, x) = \frac{1}{Z} p(y | x, \theta) p(\theta)$$

by bayes rule we have

$$\Rightarrow \theta^a = \underset{\theta}{\operatorname{argmax}} \mathcal{N}(y | x^T \theta, \sigma^2) \mathcal{N}(\theta | 0, \sigma^2 I)$$

$$\Rightarrow L(\theta) = \sum_{n=1}^N \log \left[\frac{1}{\sigma^2 \sqrt{2\pi}} e^{-\frac{(y_n - x_n^T \theta)^2}{2\sigma^2}} \right] + \sum_{n=1}^D \log \left[\frac{1}{\sigma^2 \sqrt{2\pi}} e^{-\frac{\theta_n^2}{2\sigma^2}} \right]$$

we apply log and minus and the likelihood function will be

$$= -N \log \sigma^2 \sqrt{2\pi} - \frac{1}{2\sigma^2} \sum_{n=1}^N (y_n - x_n^T \theta)^2 - D \log(\sigma^2 \sqrt{2\pi}) - \frac{1}{2\sigma^2} \sum_{n=1}^D \theta_n^2$$

$$= \underbrace{-\sum_{n=1}^N (y_i - x_i^T \theta)^2}_{\text{MSE}} - \underbrace{\frac{1}{2\sigma^2} \|\theta\|_2^2}_{\text{Regularization term with } \lambda = \frac{1}{\sigma^2}} - \underbrace{D \log(\sigma^2 \sqrt{2\pi}) - N \log \sigma^2 \sqrt{2\pi}}_{\text{constant}}$$

Hence it is equivalent to minimizing with MSE and regularization term of type L2.

$$\textcircled{C} \quad P(\theta) = \prod_{m=1}^M p(\theta_m) = \prod_{m=1}^M \frac{1}{2b} e^{-\frac{|\theta_m|}{b}} \quad (\mu=0)$$

like the previous Questions are here:

$$P(\theta | x, y) = \frac{1}{Z} P(y | x, \theta) P(\theta) = \frac{1}{Z} N(y | \theta^T x, \sigma^2) L(\theta | D, b)$$

$$\Rightarrow L(\theta) = \sum_{i=1}^N \log \frac{1}{\sigma^2 \sqrt{2\pi}} e^{-\frac{(y_i - x_i^T \theta)^2}{2\sigma^2}} + \sum_{j=1}^M \log \frac{1}{2b} e^{-\frac{|\theta_j|}{b}}$$

$$= \underbrace{\frac{1}{2\sigma^2} \sum_{i=1}^N (y_i - x_i^T \theta)^2}_{\text{MSE}} = \underbrace{\frac{1}{b} \sum_{j=1}^M \|\theta_j\|}_{\text{regularization}} - \underbrace{N \log \sigma^2 \sqrt{2\pi} - \log 2b}_{\text{constant}}$$

Hence it's like minimizing with MSE with L1 regularization
with $\lambda = \frac{2\sigma^2}{b}$

```

In [ ]: import numpy as np
import matplotlib.pyplot as plt

# Input: number of samples N
# Output: one-dimensional data set of N points where  $y_n = \sin(2 \pi x_n) + \epsilon_n$  as in lecture
def sine_data_set(N):
    np.random.seed(1234)
    x = np.random.uniform(0,1,(N))
    y = np.sin(x*2*np.pi)+np.random.normal(scale=0.2,size=(N))
    return x,y

# Input: instances X as N x M matrix, model theta
# Output: predictions of model theta on X
def predict_regression(X,theta):
    return X @ theta

# Input: one-dimensional inputs x as vector of length N, polynomial degree d
# Output: polynomial feature representation of the inputs as N x (d+1) matrix
def poly_features(x,d):
    X = np.zeros((x.shape[0],d+1))
    for i in range(0,d+1):
        X[:,i] = np.power(x,i)
    return X

```

```

In [ ]: # Generate and plot toy data set
N = 100 #3,10,20,100
x,y = sine_data_set(N)
plt.xlim([0,1])
plt.ylim([-1.2,1.2])
plt.scatter(x,y)

# Polynomial feature representation of degree d
d = 10
X = poly_features(x,d)

sigma = 0.1


```

```

sigma_p = 100
M = d + 1

# Compute MAP model
I = np.eye(M)
sigma_sqr = sigma ** -2
sigma_p_sqr = sigma_p ** -2
theta_MAP = sigma_sqr * np.linalg.inv((sigma_sqr * X.T @ X) + (sigma_p_sqr * I)) @ X.T @ y
# Plot MAP model
x_axis = np.arange(0, 1, 0.01)
x_values = poly_features(x_axis, d)
predictions = predict_regression(x_values, theta_MAP)
plt.plot(x_axis, predictions, label="predictions", c='r')
plt.show()

```



```

In [ ]: # Compute posterior distribution
I = np.eye(M)

```

```

A = (sigma_sqr * X.T @ X) + (sigma_p_sqr * I)
A_inv = np.linalg.inv(A)
theta_bar = sigma_sqr * A_inv @ X.T @ y
print(f'A is {A} - and theta_bar is {theta_bar}')

```


```

A is [[10000.0001      5183.62907879  3452.73380164  2545.0681362
      1995.32757289  1632.80140724  1378.61667748  1191.66060306
      1048.75864724   936.05510082   844.8568212 ]
 [ 5183.62907879  3452.73390164  2545.0681362   1995.32757289
      1632.80140724  1378.61667748  1191.66060306  1048.75864724
       936.05510082   844.8568212   769.48071413]
 [ 3452.73380164  2545.0681362   1995.32767289  1632.80140724
      1378.61667748  1191.66060306  1048.75864724   936.05510082
       844.8568212   769.48071413   706.07603138]
 [ 2545.0681362   1995.32757289  1632.80140724  1378.61677748
      1191.66060306  1048.75864724   936.05510082   844.8568212
       769.48071413   706.07603138   651.95015955]
 [ 1995.32757289  1632.80140724  1378.61667748  1191.66060306
      1048.75874724   936.05510082   844.8568212   769.48071413
       706.07603138   651.95015955   605.16728762]
 [ 1632.80140724  1378.61667748  1191.66060306  1048.75864724
       936.05510082   844.8569212   769.48071413   706.07603138
       651.95015955   605.16728762   564.30111942]
 [ 1378.61667748  1191.66060306  1048.75864724   936.05510082
       844.8568212   769.48071413   706.07613138   651.95015955
       605.16728762   564.30111942   528.27770302]
 [ 1191.66060306  1048.75864724   936.05510082   844.8568212
       769.48071413   706.07603138   651.95015955   605.16738762
       564.30111942   528.27770302   496.27268659]
 [ 1048.75864724   936.05510082   844.8568212   769.48071413
       706.07603138   651.95015955   605.16728762   564.30111942
       528.27780302   496.27268659   467.64240921]
 [ 936.05510082   844.8568212   769.48071413   706.07603138
       651.95015955   605.16728762   564.30111942   528.27770302
       496.27268659   467.64250921   441.87659063]
 [ 844.8568212   769.48071413   706.07603138   651.95015955
       605.16728762   564.30111942   528.27770302   496.27268659
       467.64240921   441.87659063   418.56524927]] - and theta_bar is [ 2.51848606e-02  8.09443585e+00 -9.26870087e+00
-5.02794255e+01
  8.29629116e+01  7.92012748e+00 -6.18397554e+01 -2.96464931e+01
  4.59159505e+01  6.33046024e+01 -5.73302872e+01]

```

```
In [ ]: # Draw 10 samples from posterior distribution and plot the corresponding models
n_samples = 10
samples = np.random.multivariate_normal(mean=theta_bar, cov=A_inv, size=n_samples)
plt.figure(figsize=(8,5))
for i in range(n_samples):
    plt.plot(samples[i], label=f"No. {i}")
plt.title(f'Sample of {n_samples}')
plt.legend()
plt.show()
```

C:\Users\Amir Hossein\AppData\Local\Temp\ipykernel_19400\2040710156.py:3: RuntimeWarning: covariance is not symmetric positive-semidefinite.
 samples = np.random.multivariate_normal(mean=theta_bar, cov=A_inv, size=n_samples)




```

In [ ]: N_array = [3, 10 , 100]
for N in N_array:
    x,y = sine_data_set(N)
    plt.xlim([0,1])
    plt.ylim([-1.2,1.2])
    plt.scatter(x,y)
    d = 10
    X = poly_features(x,d)


    sigma = 0.1
    sigma_p = 100
    M = d + 1


    # Compute MAP model
    I = np.eye(M)
    sigma_sqr = sigma ** -2
    sigma_p_sqr = sigma_p ** -2
    theta_MAP = sigma_sqr * np.linalg.inv((sigma_sqr * X.T @ X) + (sigma_p_sqr * I)) @ X.T @ y
    # Plot MAP model
    x_axis = np.arange(0, 1, 0.01)
    x_values = poly_features(x_axis, d)
    predictions = predict_regression(x_values, theta_MAP)
    plt.plot(x_axis, predictions, label="predictions", c='r')
    plt.show()
    A = (sigma_sqr * X.T @ X) + (sigma_p_sqr * I)
    A_inv = np.linalg.inv(A)
    theta_bar = sigma_sqr * A_inv @ X.T @ y
    n_samples = 10
    samples = np.random.multivariate_normal(mean=theta_bar, cov=A_inv, size=n_samples)
    plt.figure(figsize=(8,5))
    for i in range(n_samples):
        plt.plot(samples[i], label=f"No. {i}")
    plt.title(f'Sample of {n_samples}')
    plt.legend()
    plt.show()

```






Sample of 10






Sample of 10





```
C:\Users\Amir Hossein\AppData\Local\Temp\ipykernel_19400\1959851063.py:30: RuntimeWarning: covariance is not symmetric positive-semidefinite.  
  samples = np.random.multivariate_normal(mean=theta_bar, cov=A_inv, size=n_samples)
```



In []: