

# Linear Regression

Lecture Series „Machine Learning“

Niels Landwehr

Research Group „Data Science“  
Institute of Computer Science  
University of Hildesheim

# Agenda

- Linear regression model and normal equations
- Gradient descent for parameter estimation

# Agenda

- Linear regression model and normal equations
- Gradient descent for parameter estimation

# Review: Supervised Learning

- Review: in **supervised learning**, the goal is to make predictions about objects

Input: an **instance** (object)  $\mathbf{x}$   
that lives in an instance space  $\mathcal{X}$

$$\mathbf{x} \mapsto y$$

$\mathbf{x} \in \mathcal{X}$        $y \in \mathcal{Y}$

Output: a **label** or **target**  $y$  that  
lives in a target space  $\mathcal{Y}$

- To obtain predictions, we are looking for a **model**  $f$  that produces a prediction  $f(\mathbf{x}) \in \mathcal{Y}$  for an input instance  $\mathbf{x}$

$$f: \mathcal{X} \rightarrow \mathcal{Y}$$

Input: instance  $\mathbf{x}$

$$\mathbf{x} \mapsto f(\mathbf{x})$$

Output: prediction  $f(\mathbf{x})$

- Model will be inferred from **training data**: a set of instances with observed targets

$$\mathcal{D} = \{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)\}$$

Training instances  $\mathbf{x}_n \in \mathcal{X}$ : observed objects in  
training data, for example flowers, images of  
digits, or emails

Observed labels or targets  $y \in \mathcal{Y}$  in training  
data, for example classes of flowers, digits  
0...9, or spam/legitimate classifications

# Models as Parameterized Functions

- **How does the model  $f$  look like?**
- An important and frequently used class of models are **parameterized functions**:
  - The model, often written as  $f_{\theta}$ , is a function parameterized by a vector of parameters  $\theta \in \mathbb{R}^D$
  - The model is determined by the structure of the function and its parameters  $\theta$
  - The structure of the function is chosen a priori, the parameters are chosen during learning based on the training data
- In this setting, there are basically two ingredients to a machine learning approach:
  - We need to choose the set of models under consideration. This is determined by the structure of the function: it is given by the set

$$\mathcal{F} = \{f_{\theta} \mid \theta \in \mathbb{R}^D\}$$

which is also called the model space

- Given the model space, we need to determine how we choose a particular model  $f_{\theta^*} \in \mathcal{F}$  based on the training data. This is the actual learning algorithm.

# Linear Regression

- A simple but important instance of parameterized functions in machine learning are so-called **linear models**, for regression (this lecture) and classification (next lecture)
- **Linear Regression:** a simple but widely used model for regression tasks
  - As discussed above, assume that instances are represented by vectors

$$\mathbf{x} = \begin{pmatrix} x_1 \\ x_2 \\ \dots \\ x_M \end{pmatrix} \in \mathbb{R}^M$$

- In regression tasks, the targets are continuous, that is, we assume  $y \in \mathbb{R}$
- The parameterized function is given by

Function computes  
real-valued output  
based on instance  
 $\mathbf{x} \in \mathbb{R}^M$

$$f_{\boldsymbol{\theta}}(\mathbf{x}) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots \theta_M x_M$$


Function is **linear** in input: the output is obtained by multiplying each feature  $x_m$  in the instance  $\mathbf{x}$  by a corresponding model parameter  $\theta_m$

The model parameter vector is given by

$$\boldsymbol{\theta} = \begin{pmatrix} \theta_0 \\ \theta_1 \\ \dots \\ \theta_M \end{pmatrix} \in \mathbb{R}^{M+1}$$

# Linear Regression

- To write the linear regression model more compactly, we can replace the constant term in the regression function by an additional constant element in the input vector:
  - Assume we are including an artificial constant attribute  $x_0$  in the input vector  $\mathbf{x}$  :



- Then, we can write the linear regression compactly as a dot product:

$$\begin{aligned} f_{\boldsymbol{\theta}}(\mathbf{x}) &= \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_M x_M \\ &= \theta_0 x_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_M x_M \\ &= \mathbf{x}^T \boldsymbol{\theta} \end{aligned}$$

- In this notation, the constant term  $\theta_0$  is a model parameter like any other  $\theta_m$ , only that it happens to be in front of a constant feature

# Learning by Optimizing a Loss Function

- For linear regression, the linear regression equation defines the model space: it contains all functions that have this form, that is,

$$\mathcal{F} = \{f_{\boldsymbol{\theta}} \mid f_{\boldsymbol{\theta}}(\mathbf{x}) = \mathbf{x}^T \boldsymbol{\theta}, \boldsymbol{\theta} \in \mathbb{R}^D\}$$

- It remains to specify how we choose the model parameters  $\boldsymbol{\theta}$  (and thereby a specific model out of the model space) based on the training data  $\mathcal{D} = \{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)\}$
- How do we choose model parameters  $\boldsymbol{\theta}$  based on training data?
  1. Need to quantify how good a model parameter vector  $\boldsymbol{\theta}$  is, based on training data: this is accomplished by defining a **loss function**
  2. Need to find model parameters that are good according to loss function: this results in an **optimization problem** that has to be solved by the machine learning algorithm



# Loss Functions

- **Loss function:** measure how well a model  $f_{\theta}$  fits the training data
- Intuitively, the loss measures how close the predictions of the model on the training instances,  $f_{\theta}(\mathbf{x}_n)$ , are to the observed targets,  $y_n$ 
  - Low loss: predictions of the model are close to target and the model captures the underlying  $(\mathbf{x}, y)$ -relationship in the data well
  - High loss: predictions of the model are far away from target and model does not capture underlying  $(\mathbf{x}, y)$ -relationship in the data
  - The loss function is also sometimes called cost function or objective function
- Loss on a single training example  $(\mathbf{x}_n, y_n)$  is defined by an instance-level loss function

$$\ell(f_{\theta}(\mathbf{x}_n), y_n)$$

- Loss on entire data set is the average of the losses on all instances and defined by

Here, we make the dependence on the model parameters  $\theta$  explicit, because we want to optimize the loss in  $\theta$

$$L(\theta) = \frac{1}{N} \sum_{n=1}^N \ell(f_{\theta}(\mathbf{x}_n), y_n)$$

# Loss Functions

- Example: squared and absolute loss for regression
  - Most common loss function for regression is the **squared loss**:

$$\ell(f_{\theta}(\mathbf{x}_n), y_n) = (f_{\theta}(\mathbf{x}_n) - y_n)^2$$

Squared difference between prediction  $f_{\theta}(\mathbf{x}_n)$  and target  $y_n$ . Penalizes outliers disproportionately.

- Alternatively, the **absolute loss** can be used:

$$\ell(f_{\theta}(\mathbf{x}_n), y_n) = |f_{\theta}(\mathbf{x}_n) - y_n|$$

Absolute difference between prediction  $f_{\theta}(\mathbf{x}_n)$  and target  $y_n$ .


- Example: squared loss for linear regression  $f_{\theta}(\mathbf{x}) = 0.5\mathbf{x} + 1$  with  $\mathbf{x} \in \mathbb{R}$

One-dimensional linear regression model  $f_{\theta}(\mathbf{x}) = 0.5\mathbf{x} + 1$  given by blue line in graph.

Loss is computed at  $\mathbf{x}_n = 10$ .

Function value at this point is  $f_{\theta}(\mathbf{x}_n) = 6$ , target at this point is  $y_n = 10$

Squared loss is  $\ell(f_{\theta}(\mathbf{x}_n), y_n) = (6 - 10)^2 = 16$ .



# Learning: Minimizing the Loss Function

- The loss function defines how well a model fits the training data, and thereby how well the model captures the underlying  $(\mathbf{x}, y)$ -relationship in the data
- Given the definition of a loss function, we can view the problem of learning the model  $f_{\theta}$  from data as an optimization problem: find the model parameters with the lowest loss on the training data

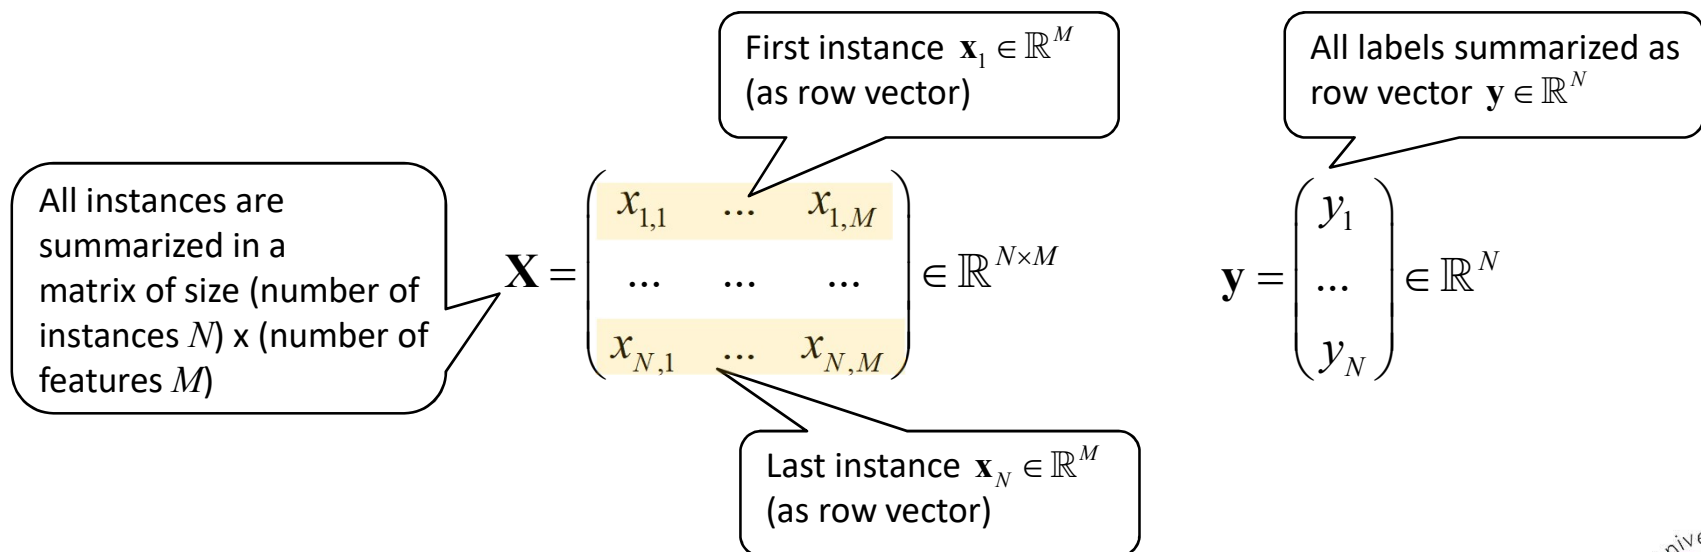
**Learning as  
optimization:**

$$\begin{aligned}\theta^* &= \arg \min_{\theta} L(\theta) \\ &= \arg \min_{\theta} \frac{1}{N} \sum_{n=1}^N \ell(f_{\theta}(\mathbf{x}_n), y_n)\end{aligned}$$

- The model  $f_{\theta^*}$  with loss-minimizing parameters  $\theta^*$  is the learned model
- **Remark:** in addition to simply minimizing the loss on the training data, in practice often need to augment the optimization criterion with a regularization term to prevent overfitting. We will come back to this in a later lecture.

# Training Data in Matrix Form

- How do we solve the optimization problem  $\boldsymbol{\theta}^* = \arg \min_{\boldsymbol{\theta}} L(\boldsymbol{\theta})$ ?
  - Approximately or exactly?
  - Computational efficiency?
  - Numerical stability?
- To simplify notation, let us write the training data  $\mathcal{D} = \{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)\}$  in matrix form:



# Training Data in Matrix Form

- Let  $\hat{y}_n = f_{\theta}(\mathbf{x}_n)$  denote the prediction of the model on the  $n$ -th instance
- The predictions of a linear model  $f_{\theta}(\mathbf{x}) = \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_M x_M$  on the entire training data can now be written as

$$\hat{\mathbf{y}} = \mathbf{X}\boldsymbol{\theta} \in \mathbb{R}^N$$

Result of multiplying  $N \times M$  matrix  $\mathbf{X}$  with  $M$ -dimensional vector  $\boldsymbol{\theta}$  is  $N$ -dimensional vector  $\hat{\mathbf{y}}$ . By construction of the matrix  $\mathbf{X}$ ,  $n$ -th element of result vector is  $\theta_1 x_{n,1} + \theta_2 x_{n,2} + \dots + \theta_M x_{n,M}$

- Here, we have left out the constant model term  $\theta_0$  because it can be incorporated into an additional constant feature as described above
- Define the L2-norm of a vector  $\mathbf{z} \in \mathbb{R}^N$  as

$$\|\mathbf{z}\|_2 = \sqrt{z_1^2 + \dots + z_N^2}$$

- The squared loss of the linear model  $f_{\theta}(\mathbf{x})$  can be computed as

$$L(\boldsymbol{\theta}) = \frac{1}{N} \|\mathbf{y} - \mathbf{X}\boldsymbol{\theta}\|_2^2$$

Here, the squared L2-norm sums up the squared differences between the targets in  $\mathbf{y}$  and the predictions in  $\mathbf{X}\boldsymbol{\theta}$

# Normal Equations for Linear Regression

- The model parameters minimizing the squared loss,

$$\begin{aligned}\boldsymbol{\theta}^* &= \arg \min_{\boldsymbol{\theta}} \frac{1}{N} \|\mathbf{y} - \mathbf{X}\boldsymbol{\theta}\|_2^2 \\ &= \arg \min_{\boldsymbol{\theta}} \|\mathbf{y} - \mathbf{X}\boldsymbol{\theta}\|_2^2\end{aligned}$$

constant 1/N does not affect minimum

are also called the least squares estimates, because they minimize the squared errors

- The least squares estimates can be computed by solving for  $\boldsymbol{\theta}$  the so-called normal equations:

$$\mathbf{X}^T \mathbf{X} \boldsymbol{\theta} = \mathbf{X}^T \mathbf{y}$$

- Solving the normal equations means solving a system of  $M$  linear equations: solve

$$\mathbf{A} \boldsymbol{\theta} = \mathbf{b}$$

with  $\mathbf{A} = \mathbf{X}^T \mathbf{X} \in \mathbb{R}^{M \times M}$  and  $\mathbf{b} = \mathbf{X}^T \mathbf{y} \in \mathbb{R}^M$

- Different algorithmic methods available: e.g. Gaussian elimination, Cholesky decomposition, QR decomposition

# Normal Equations for Linear Regression

- Why does solving the normal equations yield the least squares model parameters?
- For two vectors  $\mathbf{a}, \mathbf{b} \in \mathbb{R}^N$  let  $\langle \mathbf{a}, \mathbf{b} \rangle = \sum_{n=1}^N a_n b_n$  denote their dot product
- Note that  $\|\mathbf{a}\|_2^2 = \langle \mathbf{a}, \mathbf{a} \rangle$ , namely the sum of squared vector elements
- The squared loss of the model  $f_{\boldsymbol{\theta}}(\mathbf{x}_n)$  can be written as

$$\frac{1}{N} \|\mathbf{y} - \mathbf{X}\boldsymbol{\theta}\|_2^2 = \frac{1}{N} \langle \mathbf{y} - \mathbf{X}\boldsymbol{\theta}, \mathbf{y} - \mathbf{X}\boldsymbol{\theta} \rangle$$

Like the squared norm, the dot product also sums up the the squared vector elements

- To find the minimum, we can set the derivative to zero

Derived by chain rule: we have an expression of the form  $\langle \mathbf{z}, \mathbf{z} \rangle$ , with  $\mathbf{z} = \mathbf{y} - \mathbf{X}\boldsymbol{\theta}$ . The outer derivative is  $\frac{\partial}{\partial \mathbf{z}} \langle \mathbf{z}, \mathbf{z} \rangle = 2\mathbf{z} = 2(\mathbf{y} - \mathbf{X}\boldsymbol{\theta})$ , similarly as in scalar calculus  $\frac{\partial}{\partial x} x^2 = 2x$ . The inner derivative is  $\frac{\partial}{\partial \boldsymbol{\theta}} \mathbf{z} = -\mathbf{X}$ . The two are multiplied resulting in  $(-\mathbf{X})^T 2(\mathbf{y} - \mathbf{X}\boldsymbol{\theta})$ .

$$\frac{\partial}{\partial \boldsymbol{\theta}} \frac{1}{N} \langle \mathbf{y} - \mathbf{X}\boldsymbol{\theta}, \mathbf{y} - \mathbf{X}\boldsymbol{\theta} \rangle = \frac{1}{N} (-\mathbf{X})^T 2(\mathbf{y} - \mathbf{X}\boldsymbol{\theta}) = -\frac{2}{N} (\mathbf{X}^T \mathbf{y} - \mathbf{X}^T \mathbf{X}\boldsymbol{\theta}) \stackrel{!}{=} 0$$

Resolving product

Set to zero to find minimum

$$\Rightarrow \mathbf{X}^T \mathbf{y} = \mathbf{X}^T \mathbf{X}\boldsymbol{\theta}$$

results in normal equations

# Normal Equations for Linear Regression

- Learning a linear regression model using normal equations, formulated as algorithm:

**Algorithm** learn-linreg-NormEq

**Input** : training data  $\mathcal{D} := \{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)\}$

**Output** : learned model parameters  $\boldsymbol{\theta}$

1.  $\mathbf{X} = (\mathbf{x}_1, \dots, \mathbf{x}_N)^T$

2.  $\mathbf{y} = (y_1, \dots, y_N)^T$

Summarize training instances in matrix  $\mathbf{X}$   
and training labels in vector  $\mathbf{y}$

3.  $\mathbf{A} = \mathbf{X}^T \mathbf{X}$

4.  $\mathbf{b} = \mathbf{X}^T \mathbf{y}$

Compute expressions  $\mathbf{A}$ ,  $\mathbf{b}$  for normal equations

5.  $\boldsymbol{\theta} = \text{solve-linear-equations}(\mathbf{A}, \mathbf{b})$

6. return  $\boldsymbol{\theta}$

Solving system of equations (using one of the  
methods mentioned above) yields final model  
parameters



# Computational Complexity

- What is the computational complexity of learning a linear regression model via the normal equations?
- Dimensions of input:

Matrix  $\mathbf{X} \in \mathbb{R}^{N \times M}$  ( $N$  = number of instances,  $M$  = number of features)

Vector  $\mathbf{y} \in \mathbb{R}^N$

- Main computational steps:

Matrix product  $\mathbf{X}^T \mathbf{X}$ :  $M \times N$  times  $N \times M$  matrix multiplication,  $O(NM^2)$

Solve system of  $M$  linear equations:  $O(M^3)$

- Overall runtime:  $O(NM^2 + M^3)$

Linear dependence on  $N$  is ok,

but computation time rises quickly with  $M$

$M$	runtime [s]
100	0.002
200	0.004
400	0.022
800	0.086
1600	0.555
3200	6.275
6400	27.902

(Intel i5-760 2.8 MHz,  
2010, Python numpy)

# Agenda

- Linear regression model and normal equations
- Gradient descent for parameter estimation

# Motivation: Gradient Descent

- For the linear regression model, the parameter vector minimizing the squared loss can be obtained by simply solving a system of linear equations
- However, this is due to the particularly simple form of the model and the particular loss function (squared loss)
- In general, for most combinations of models and loss functions, finding the model parameters that minimize the loss function is less straightforward
- Even for linear regression, computational complexity might prevent us from using normal equations to solve for the parameters minimizing the loss
- **Gradient descent:** An alternative approach for finding model parameters that (approximately) minimize the loss
  - Iterative optimization algorithm
  - Finds global optimum if the problem of loss-minimization is convex in parameters
  - Otherwise, finds a local optimum (still good & useful in many practical settings)
  - Very flexible for different models and losses, widely used in practice

# Learning: How to Solve the Optimization Problem?

- Formally speaking, to find the loss-minimizing parameter vector  $\theta$  we need to solve an optimization problem
  - Given a model class  $\mathcal{F} = \{f_{\theta} \mid \theta \in \mathbb{R}^D\}$
  - Given a loss function  $\ell(f(\mathbf{x}), y)$
  - Given training data  $\mathcal{D} = \{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)\}$
  - Find

$$\begin{aligned}\theta^* &= \arg \min_{\theta} L(\theta) \\ L(\theta) &= \frac{1}{N} \sum_{n=1}^N \ell(f_{\theta}(\mathbf{x}_n), y_n)\end{aligned}$$

- Challenges:
  - Optimization variable  $\theta$  can be high-dimensional (depending on model class)
  - Data set can be large (in terms of number of instances and/or features)
  - $L(\theta)$  can be a complex, generally non-convex function in  $\theta$
- **How to solve optimization problem?**


# Optimization: Gradient Descent

- Gradient descent: find the model parameter vector  $\theta^*$  by „walking“ in the space  $\mathbb{R}^D$  of model parameters always in the direction of steepest descent



# Loss Function Surface

- More formally: the loss function defines a surface over the  $D$ -dimensional  $\theta$ -space
- The loss function  $L(\theta)$  maps parameter vectors to a real number:  $L: \mathbb{R}^D \rightarrow \mathbb{R}$
- Can be visualized as a loss surface / landscape:



Want to find the „lowest point“:

$$\theta^* = \arg \min_{\theta} L(\theta)$$

- Note: parameter space is high-dimensional (often millions of parameters and therefore dimensions). 2D-visualizations can be misleading.
- Shape of loss surface will depend on loss function, structure of model, data



# Gradient of Function

- Idea: follow the slope. How do we get the local slope of the loss function?
- **Gradient:**

Let  $f(\mathbf{x})$  with  $f: \mathbb{R}^D \rightarrow \mathbb{R}$  denote a scalar-valued differentiable function. The gradient of  $f$ , written  $\nabla f$ , is a function  $\nabla f: \mathbb{R}^D \rightarrow \mathbb{R}^D$  defined by

$$\nabla f(\mathbf{x}) = \begin{pmatrix} \frac{\partial f}{\partial x_1}(\mathbf{x}) \\ \frac{\partial f}{\partial x_2}(\mathbf{x}) \\ \dots \\ \frac{\partial f}{\partial x_D}(\mathbf{x}) \end{pmatrix}$$

partial derivative of  $f$  with respect to  $x_1$

partial derivative of  $f$  with respect to  $x_D$

# Gradient of Function

- Idea: follow the slope. How do we get the local slope of the loss function?
- **Example for gradient (k=3):**

$$f : \mathbb{R}^3 \rightarrow \mathbb{R} \quad \text{with} \quad f(\mathbf{x}) = 3x_1^2x_2 + 2x_2^2 - x_1x_3 \quad \mathbf{x} = (x_1, x_2, x_3)^T$$

Partial derivatives:

$$\frac{\partial f}{\partial x_1} = 6x_1x_2 - x_3$$

$$\frac{\partial f}{\partial x_2} = 3x_1^2 + 4x_2$$

$$\frac{\partial f}{\partial x_3} = -x_1$$


Gradient at  $\mathbf{x} = (1, 2, 3)^T$ :

$$\nabla f(\mathbf{x}) = \begin{pmatrix} \frac{\partial f}{\partial x_1}(\mathbf{x}) \\ \frac{\partial f}{\partial x_2}(\mathbf{x}) \\ \frac{\partial f}{\partial x_3}(\mathbf{x}) \end{pmatrix} = \begin{pmatrix} 12 - 3 \\ 3 + 8 \\ -1 \end{pmatrix} = \begin{pmatrix} 9 \\ 11 \\ -1 \end{pmatrix}$$



# Gradient of Loss Function

- Idea: follow the slope. How do we get the local slope of the loss function?
- **Negative gradient points into the direction of steepest descent**




- For  $\theta \in \mathbb{R}^D$ , the vector  $-\nabla L(\theta) \in \mathbb{R}^D$  points into the direction of steepest descent
- Locally moving into the direction of negative gradient will decrease loss

$$L(\theta - \eta \nabla L(\theta)) \leq L(\theta) \quad \text{for small enough } \eta > 0$$

# Gradient Descent Algorithm (Fixed Step Size)

- Idea: follow the slope. How do we get the local slope of the loss function?
- **Gradient descent: iterative small steps in direction of negative gradient**
  - Step size: how far to move along direction of gradient? Parameter  $\eta$
  - For small enough  $\eta$ , will converge to local optimum



## Gradient descent algorithm

1.  $\theta_0 = \text{randomInitialization}()$
2. for  $i = 0, \dots, i_{\max}$ :
3.      $\theta_{i+1} = \theta_i - \eta \nabla L(\theta_i)$
4.     if  $L(\theta_i) - L(\theta_{i+1}) < \epsilon$ :
5.         return  $\theta_{i+1}$
6. raise Exception("Not converged in  $i_{\max}$  iterations")

# Computing the Gradient

- To carry out the gradient descent algorithm, we need to compute the gradient  $\nabla L(\boldsymbol{\theta})$  for any given  $\boldsymbol{\theta} \in \mathbb{R}^D$
- The function  $L(\boldsymbol{\theta})$  can be quite complex
  - parameter vector  $\boldsymbol{\theta} \in \mathbb{R}^D$  can be high-dimensional
  - computation of  $L(\boldsymbol{\theta})$  involves computing model predictions, the loss function, and a sum over data instances

$$L(\boldsymbol{\theta}) = \frac{1}{N} \sum_{n=1}^N \ell(f_{\boldsymbol{\theta}}(\mathbf{x}_n), y_n)$$

- How can we compute the gradient in practice?

# Numerical Gradients

- **First idea:** numerical gradient based on difference quotient

Let  $\boldsymbol{\theta} = (\theta_1, \dots, \theta_D)^T \in \mathbb{R}^D$ , that is,  $L : \mathbb{R}^D \rightarrow \mathbb{R}$ .

$$\text{Gradient is } \nabla L(\boldsymbol{\theta}) = \left( \frac{\partial L}{\partial \theta_1}(\boldsymbol{\theta}), \dots, \frac{\partial L}{\partial \theta_D}(\boldsymbol{\theta}) \right)^T$$

Partial derivative is limes of difference quotient:

$$\frac{\partial L}{\partial \theta_d}(\boldsymbol{\theta}) = \lim_{h \rightarrow 0} \frac{L(\boldsymbol{\theta} + h\mathbf{u}_d) - L(\boldsymbol{\theta})}{h} \quad \mathbf{u}_d = (0, 0, \dots, 0, 1, 0, \dots, 0)^T \in \mathbb{R}^D$$

$\uparrow$  d-th position

Approximate d-th entry in gradient by

$$\frac{\partial L}{\partial \theta_d}(\boldsymbol{\theta}) \approx \frac{L(\boldsymbol{\theta} + h\mathbf{u}_d) - L(\boldsymbol{\theta})}{h} \quad \text{where } h \text{ is a small number (e.g. } h = 10^{-4}\text{)}$$

# Numerical Gradients

- **First idea:** numerical gradient based on difference quotient
- To compute entire gradient, have to compute the approximation

$$\frac{\partial L}{\partial \theta_d}(\boldsymbol{\theta}) \approx \frac{L(\boldsymbol{\theta} + h\mathbf{u}_d) - L(\boldsymbol{\theta})}{h} \quad \text{for each } d \in \{1, \dots, D\}$$

- As number of parameters  $D$  can be large this is usually not efficient: Would have to compute the entire loss many times to get a single gradient
- Also, solution is only approximate and not always numerically stable
- However, numerical gradients are easy to implement and can be used to verify (debug) other, more efficient implementations of gradient computations

# Analytic Gradients

- **Second idea:** analytically derive gradient

Example from above:


$$f : \mathbb{R}^3 \rightarrow \mathbb{R} \quad \text{with} \quad f(\mathbf{x}) = 3x_1^2x_2 + 2x_2^2 - x_1x_3$$

$$\frac{\partial f}{\partial x_1} = 6x_1x_2 - x_3 \qquad \frac{\partial f}{\partial x_2} = 3x_1^2 + 4x_2 \qquad \frac{\partial f}{\partial x_3} = -x_1$$

- Advantages: exact solution, potentially faster
- This is a widely used approach for model classes and loss functions that are not too complex, such that corresponding closed-form solutions for the gradient can still be derived
- Will see examples below and in further lectures

# Automatic Differentiation


- **Alternative for models for which deriving analytic gradients is not (easily) possible: automatic differentiation**
  - algorithmic approach to compute gradient vector given model, loss, data
  - look at the overall expression defining  $L(\theta)$  as a graph of elementary operations („computation graph“ or „data flow graph“)
  - if we know the derivatives of the individual operations, we can efficiently compute the overall derivative by the chain rule
  - widely used for neural networks



- No details here, will discuss this in more detail when talking about neural networks

# Step Sizes in Gradient Descent

- For the gradient descent algorithm, we need to specify a step size  $\eta$
- A step size is needed, because the negative gradient  $-\nabla L(\theta)$  only points into the direction of steepest descent in an infinitesimal neighborhood of  $\theta$
- Therefore, a decrease in loss is only guaranteed for small enough  $\eta$




- If step size is too large, algorithm might not find any good  $\theta$
- If step size is too small, algorithm will take very long to find good  $\theta$





# Selecting Step Sizes in Gradient Descent

- How to select step size?
- Option 1: fixed step size which is a tunable hyperparameter (that is, have to empirically test different step sizes and pick the one that works best)
- Option 2: policy to reduce step size over time or adapt it heuristically (e.g. by monitoring how much parameters change at each iteration)
- Option 3: rather than using a fixed step size, perform a **line search** along the direction of the negative gradient



# Local Minima


- In general, the function  $L(\theta)$  can have several different **local minima**, that is, parameters  $\theta$  with  $\nabla L(\theta) = 0$
- Gradient descent will find one of the local minima, which is not necessarily globally optimal
- The gradient can also become zero at a saddle point, and (simple) gradient descent can then get „stuck“



# Convexity

- A function  $L: \mathbb{R}^D \rightarrow \mathbb{R}$  is called **convex** if for any  $\theta_1, \theta_2 \in \mathbb{R}^D$  and any  $t \in [0, 1]$

$$L(t\theta_1 + (1-t)\theta_2) \leq tL(\theta_1) + (1-t)L(\theta_2)$$



- For a convex function, any local minimum is also a global minimum
- For a convex function, gradient descent will thus find a global optimum, given sufficiently small step sizes

# Convexity

- If the function  $L : \mathbb{R}^D \rightarrow \mathbb{R}$  is twice differentiable, it is convex if the Hessian matrix

$$\mathbf{H} = \begin{pmatrix} \frac{\partial L}{\partial \theta_1 \partial \theta_1} & \cdots & \frac{\partial L}{\partial \theta_1 \partial \theta_D} \\ \vdots & \ddots & \vdots \\ \frac{\partial L}{\partial \theta_D \partial \theta_1} & \cdots & \frac{\partial L}{\partial \theta_D \partial \theta_D} \end{pmatrix}$$

is positive semidefinite for all  $\boldsymbol{\theta} \in \mathbb{R}^D$  (remember that  $\mathbf{H}$  depends on  $\boldsymbol{\theta}$  )

- The matrix  $\mathbf{H}$  is positive semidefinite if for all  $\bar{\boldsymbol{\theta}} \in \mathbb{R}^D$  it holds that  $\bar{\boldsymbol{\theta}}^T \mathbf{H} \bar{\boldsymbol{\theta}} \geq 0$
- For any matrix  $\mathbf{A} \in \mathbb{R}^{N \times M}$ , the matrix  $\mathbf{A}^T \mathbf{A}$  is positive semidefinite
- For linear regression with squared loss, the Hessian of the loss of the model in the model parameters is  $(2 / N) \mathbf{X}^T \mathbf{X}$  (can be derived similarly as for normal equations)
  - Loss is therefore convex and gradient descent will find global optimum

# Learning Linear Regression by Gradient Descent

- Instead of solving the normal equations, we can also use gradient descent to learn a linear regression model
- For linear regression, the gradient of the loss function is (result copied from derivation of normal equations above):

$$\nabla L(\boldsymbol{\theta}) = \frac{\partial}{\partial \boldsymbol{\theta}} \frac{1}{N} \|\mathbf{X}\boldsymbol{\theta} - \mathbf{y}\|_2^2 = \frac{1}{N} (-\mathbf{X})^T 2(\mathbf{y} - \mathbf{X}\boldsymbol{\theta})$$

- Can be computed readily from training data and current parameter vector
- Pick learning rate  $\eta$  and run gradient descent algorithm:

## Gradient descent algorithm

1.  $\boldsymbol{\theta}_0 = \text{randomInitialization}()$
2. for  $i = 0, \dots, i_{\max}$ :
3.      $\boldsymbol{\theta}_{i+1} = \boldsymbol{\theta}_i - \eta \nabla L(\boldsymbol{\theta}_i)$
4.     if  $L(\boldsymbol{\theta}_i) - L(\boldsymbol{\theta}_{i+1}) < \epsilon$ :
5.         return  $\boldsymbol{\theta}_{i+1}$
6. raise Exception("Not converged in  $i_{\max}$  iterations")

# Sparsity in Linear Regression

- In some application domains, instance vectors  $\mathbf{x} \in \mathbb{R}^M$  will be high-dimensional but sparse, meaning that most of the entries in a vector are typically zero
- For example, using a word indicator representation for spam classification: there are tens of thousands of words, but only a few of them appear in the average message
- In gradient descent, we can take advantage of this sparsity:

Multiplication of sparse matrix with dense vector is efficient: only go over non-zero entries in matrix

$$\nabla L(\boldsymbol{\theta}) = \frac{1}{N} (-\mathbf{X})^T 2(\mathbf{y} - \mathbf{X}\boldsymbol{\theta})$$

Again, multiplication of sparse matrix with dense vector

- Computational complexity of gradient descent is then  $O(I(NM_{nz} + M))$ 
  - $I$  is number of iterations,  $M_{nz}$  average number of nonzero elements in row of  $\mathbf{X}$
- In contrast, approaches to solving the normal equations cannot easily take advantage of sparsity

# Example: Linear Regression

- **Example: gradient descent for linear regression**

One-dimensional model  $f: \mathbb{R} \rightarrow \mathbb{R}$ , using a constant attribute  $x_2 = 1$ :

$$f_{\theta}(\mathbf{x}) = \theta_1 x_1 + \theta_2 x_2 = \theta_1 x_1 + \theta_2$$

↑ input


Assume three training instances. The training data can be given in matrix form with  $\mathbf{X} \in \mathbb{R}^{3 \times 2}$  and  $\mathbf{y} \in \mathbb{R}^3$ :

$$\mathbf{X} = \begin{pmatrix} 1 & 1 \\ 2 & 1 \\ 3 & 1 \end{pmatrix}$$

↑      ↑  
inputs  $x_1$       constant  
                         attribute  $x_2$

$$\mathbf{y} = \begin{pmatrix} 3 \\ 4 \\ 6 \end{pmatrix}$$

↑  
targets  $y$



# Example: Linear Regression

- Example: gradient descent for linear regression

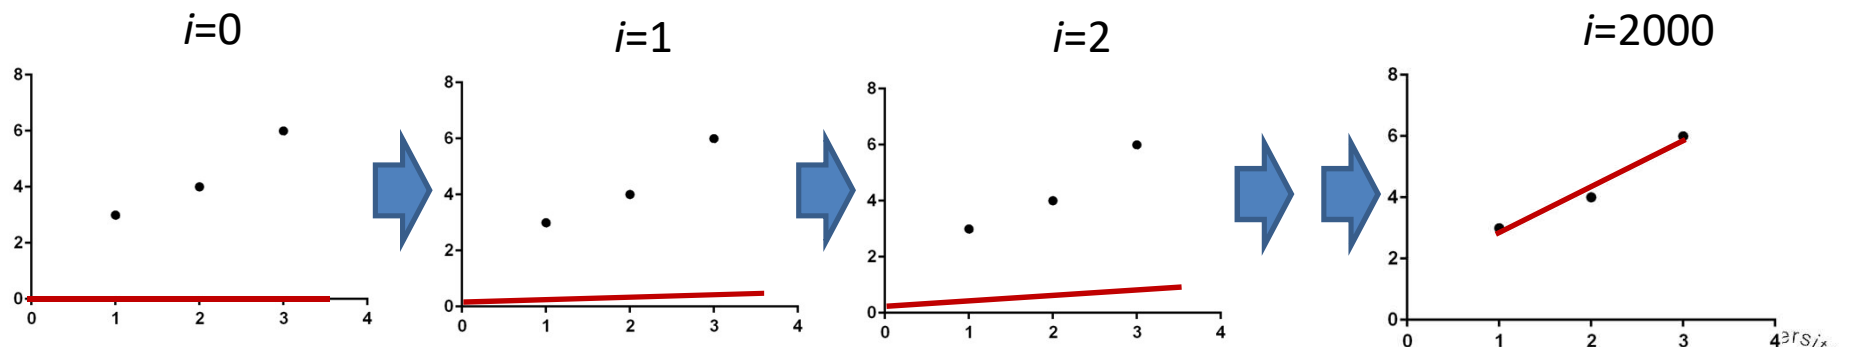
We can train the model using a squared loss function and gradient descent:

$$\theta^* = \arg \min_{\theta} L(\theta)$$
$$L(\theta) = \frac{1}{3} \sum_{n=1}^3 (f_{\theta}(\mathbf{x}_n) - y_n)^2$$

## Gradient descent algorithm

1.  $\theta_0 = \text{randomInitialization}()$
2. for  $i = 0, \dots, i_{\max}$ :
3.    $\theta_{i+1} = \theta_i - \eta \nabla L(\theta_i)$
4.   if  $L(\theta_i) - L(\theta_{i+1}) < \epsilon$ :
5.     return  $\theta_{i+1}$
3. raise Exception("Not converged in  $i_{\max}$  iterations")

For example, initialize to zeros






# Gradient Descent: Example

- Example: gradient descent for linear regression


Development of model parameters  $\theta_1, \theta_2$  during gradient descent



# Gradient Descent: Example

- **Example: gradient descent for linear regression**

Development of loss function during gradient descent



# Summary

- **Linear regression** is a simple but widely used model for regression problems
- The model can be learned from data by minimizing a **loss function**, e.g. squared loss
- Minimizing the loss function is an optimization problem, which for linear regression can be solved by solving the normal equations (system of linear equations)
- Alternatively, **gradient descent** is a flexible iterative optimization method that can be used to learn linear regression and also many other models
  - start with random parameter vector
  - repeatedly update the parameter vector by following the direction of steepest descent, given by the negative gradient